



ugr

Universidad
de Granada

PROYECTO DE FIN DE CARRERA
INGENIERÍA EN INFORMÁTICA

GNU Psychosynth

A framework for modular, interactive and
collaborative sound synthesis and live music
performance

Autor

Juan Pedro Bolívar Puente

Director

Joaquín Fernández-Valdivia



DECSAI
Universidad de Granada

DEPARTAMENTO DE CIENCIAS DE LA COMPUTACIÓN E
INTELIGENCIA ARTIFICIAL

—
Granada, Junio de 2011



GNU Psychosynth

**A framework for modular, interactive and
collaborative sound synthesis and live music
performance**

Autor

Juan Pedro Bolívar Puente

Directores

Joaquín Fernández-Valdivia

GNU Psychosynth: A framework for modular, interactive and collaborative sound synthesis and live music performance

Juan Pedro Bolívar Puente

Palabras clave:

Resumen

Escribir el abstract en español.

GNU Psychosynth: A framework for modular, interactive and collaborative sound synthesis and live music performance

Juan Pedro Bolívar Puente

Keywords:

Abstract

Escribir el abstract en inglés.

Yo, **Juan Pedro Bolívar Puente**, alumno de la titulación Ingeniería en Informática de la **Escuela Técnica Superior de Ingenierías Informática y de Telecomunicación de la Universidad de Granada**, con DNI 48941569F, autorizo la ubicación de la siguiente copia de mi Proyecto Fin de Carrera en la biblioteca del centro para que pueda ser consultada por las personas que lo deseen.

Fdo: Juan Pedro Bolívar Puente

Granada a 1 de Junio de 2011.

D. **Joaquín Fernández-Valdivia**, Catedrático del Departamento de Ciencias de la Computación e Inteligencia Artificial de la Universidad de Granada.

Informan:

Que el presente proyecto, titulado *GNU Psychosynth: A framework for modular, interactive and collaborative sound synthesis and live music performance*, ha sido realizado bajo su supervisión por **Juan Pedro Bolívar Puente**, y autorizamos la defensa de dicho proyecto ante el tribunal que corresponda.

Y para que conste, expiden y firman el presente informe en Granada a 1 de Junio de 2010.

El director:

D. **Joaquín Fernández-Valdivia**

AGRADECIMIENTOS

Añadir agradecimientos.

Sinceramente gracias.

CONTENTS

Preface — A personal, historical and audiophile dissertation	1
1 Introduction, definition and goals	9
1.1 Problem definition	9
1.1.1 A modular synthesiser	10
1.1.2 An interactive synthesiser	13
1.1.3 A collaborative environment	16
1.1.4 A framework	19
1.2 Objectives	20
1.2.1 Main objective	20
1.2.2 Preconditional objectives	21
1.2.3 Partial objectives	22
1.3 Background and state of the art	23
1.3.1 Modular synthesis	23
1.3.2 Touchable and tangible interfaces	26
1.3.3 Audio synthesis libraries and frameworks	29
2 Analysis and planning	31
2.1 Requirement modelling	31
2.1.1 Functional requirements	32
2.1.2 Non-functional requirements	39
2.2 Open, community oriented development	40
2.2.1 Software License	41
2.2.2 The GNU project	41
2.2.3 Software forge and source code repository	41

2.2.4	Communication mechanisms	42
2.3	Development environment	43
2.3.1	Programming language	44
2.3.2	Operating System	45
2.3.3	Third party libraries	46
2.4	Architecture and current status	49
2.4.1	A critique on the term framework	51
2.4.2	The layered architecture	51
2.5	Project planning and methodology	58
2.5.1	Rationale — A critique on software engineering	58
2.5.2	An iterative development model	60
2.5.3	A project plan	63
3	A generic sound processing library	67
3.1	Analysis	69
3.1.1	Factors of sound representation	69
3.1.2	Common solutions	74
3.1.3	A generic approach: Boost.GIL	75
3.2	Design	77
3.2.1	Core techniques	77
3.2.2	Core concepts	85
3.2.3	Going dynamic	90
3.2.4	Input and Output module	90
3.2.5	Synthesis module	90
3.3	Validation	90
3.3.1	Unit testing	90
3.3.2	Performance	90
3.3.3	Integration	90
3.4	Conclusions and future work	92
	Bibliography	99

Appendices	103
A Glossary	103
B Gnu General Public License	105
B.1 Preamble	105
B.2 Terms and Conditions	107

LIST OF FIGURES

Figure 0.1	Keith Emerson playing a modular Moog	3
Figure 0.2	Artwork for the music of Kraftwerk	7
Figure 0.3	Autechre related Max/MSP patches	8
Figure 1.1	A screenshot of Psychosynth 0.1.1 using dynamic patching to connect a complex graph.	15
Figure 1.2	The JazzMutant's Lemur touchable music interface.	16
Figure 1.3	A conga-alike MIDI controller.	17
Figure 1.4	An example of the Reactable being used collaboratively by several people.	18
Figure 1.5	An example of Psychosynth 0.1.4 being used collaboratively over the network.	19
Figure 1.6	Native Instrument's Reaktor editing a <i>core cell</i> based patch.	25
Figure 1.7	A typical Reactable setup.	27
Figure 1.8	Electronic music artist Four Tet performing with a Tenori-On	28
Figure 2.1	A screenshot of GNU Psychosynth 0.1.4	50
Figure 2.2	The Psychosynth layered architecture.	52
Figure 2.3	Representation of the node graph as in Psychosynth 0.1.7	55
Figure 2.4	Communication between the audio and user interface thread as in Psychosynth 0.1.7	55

List of Figures

Figure 2.5	The MVC architectural style	56
Figure 3.1	Multi-channel data in planar (a) and interleaved (b) form.	74

LIST OF TABLES

Table 3.1	Terminology map from <code>boost::gil</code> to <code>psynth::sound</code>	76
-----------	---	----

TODO LIST

Escribir el abstract en español.	5
Escribir el abstract en inglés.	7
Añadir agradecimientos.	13
En el resto del documento uso el plural de modes- tia, pero aquí hablo en un tono más informal y en primera persona. La idea es introducir al lector, que probablemente esté totalmente desfa- miliarizado con el tema, haciendo un recorrido cronológico en paralelo de la concepción y desar- rollo del Psychosynth así como de mi interés por la música electrónica y de la música electrónica en sí (al fin y al cabo, mi interés por la música electrónica está ordenado por la su cronología también.).	1
Al final lo he pasado al prefacio. Si alguien cree mejor ponerlo en la introducción, pues que lo diga. . . .	1
Me consta que David ha trabajado en la industria de videojuegos sintetizando sonidos pero debería preguntarle a ver si hay algo más destacable en su currículum.	21
A continuación hablo de aplicaciones comerciales que no tienen (o no he encontrado) artículos academi- cos. Estoy evitando citar los manuales de usuario o simplemente enlazar a sus webs comerciales, pero no sé que es lo más ortodoxo hacer en estos casos. En algunos casos más adelante y cuando lo he considerado relevante estoy poniendo en- laces al pie de página cuando no existe un paper académico. Cualquier recomendación al respecto es bienvenida.	24

De hecho, la naturaleza real de las señales de control aún no la tengo clara al 100%. Las señales de control son lo que en la versión actual de GNU/Psychosynth se llaman <i>parámetros</i> . Sin embargo, el tema no es del todo trivial y tras ver como funciona Reaktor y pensar mucho en como reinventar el MVC del Psychosynth la situación está cada vez más confusa, pero supongo que la concreción se puede dejar para el rediseño de esta capa del software en la segunda iteración. . .	33
En realidad este requisito es un poco ambigüo, pero no se muy bien cómo especificar la posibilidad e usar tamaños de buffer bajos, ya que las condiciones en las que es medible esto dependen de la potencia del procesador y, en un sinte modular que permite redes arbitrariamente complejas, del tamaño del sinte. ¿Alguna sugerencia?	40
Aquí estoy citando un blog — ¿Es esto correcto o mejor lo pongo como nota al pie?	44
Mejorar y traducir esta figura.	56
Lo se, soy un criticón enfermizo. He acabado convirtiendo esta sección en una crítica (indirecta y relativamente elegante, espero) a la concepción de la ingeniería del software que se da en la facultad y que se hace a menudo de los PFC. Si pensáis que esta crítica es demasiado larga o inapropiada puedo matizarla o eliminarla :D No obstante, en su defensa la considero apropiada ya que la elección de una metodología debe ser racional y como tanto debe incluir un análisis crítico de las alternativas.	58
Este capítulo en general necesita más dibujos.	69

Me he dado cuenta de que esta sección introduce
muchas definiciones. ¿Tal vez utilizar un estilo
más formal usando el paquete theorem para las
definiciones ayudaría a usar el documento como
referencia? 69

¿Añado un glosario recopilando las definiciones que
doy de diferentes términos específicos del do-
minio del audio y para expandir y aclarar siglas
y ac[Pleaseinsertintopreamble]onimos? ¿O mejor
paso? — JP 103

PREFACE — A PERSONAL, HISTORICAL AND AUDIOPHILE DISSERTATION

En el resto del documento uso el plural de modestia, pero aquí hablo en un tono más informal y en primera persona. La idea es introducir al lector, que probablemente esté totalmente desfasado con el tema, haciendo un recorrido cronológico en paralelo de la concepción y desarrollo del Psychosynth así como de mi interés por la música electrónica y de la música electrónica en sí (al fin y al cabo, mi interés por la música electrónica está ordenado por la su cronología también.).

Al final lo he pasado al prefacio. Si alguien cree mejor ponerlo en la introducción, pues que lo diga.

I am going to let the formalities, both in form and content, of a final project aside in this section to give an initial background on the historical development of the conception of GNU Psychosynth. After all, the story of this project is, in many ways, the story of my own learning and maturing process and specially the evolution of my interest in music.

In 2006 I was a young computer science student who had just moved to Granada, a city full of youth, joy and cultural activities. At that time, I was not keen at all in electronic music —maybe biased by my prejudices on the rave subculture that surrounds a wide part of it, even though eventually I happened to appreciate it in some way. At that time I was

more of a punk and ska fan and rejected the artificiality and production process of electronic music; this was actually a contradiction with my interest in programming. However, I eventually got specially interested in the wild 70's explosion of musical experimentation, and concretely in progressive rock.

I can vividly remember my first positive contact with electronic music. It was a chilled and psychedelic evening at a friend's place —one of those old and rundown but magical flat, with rather high ceilings and wooden windows, where many students live in Granada— when we Youtubed a video where Keith Emerson virtuously performed “Knife Edge” on a keyboard attached to a big wooden box full of wires and knobs. By rearranging the wires or playing with the knobs, he would radically change the sound that his keyboard emitted. That magical sound machine was a Moog modular synthesiser, and that was the birth of a true love for electronic music that would later conceive Psychosynth —is not love, they say, the true means for conception?

After that, I started to listen to more and more music made with electronic devices — an exploration that happened, actually, following electronic music's own history. From the synthesiser-full rock of Soft Machine, King Crimson or The Nice I opened my ears to the purely synthetic orchestral compositions of Wendy Carlos, Vangelis or the early Jean Michelle Jarre. Kraftwerk's own evolution from rock to drum-pads, vocoders and synthesisers allowed me to open my ears to more modern electronic music.

It was still my first year of university when I watched a video, under circumstances probably similar to that before, of a new device being developed in the University Pompeu Fabra, in Barcelona: the ReacTable [1]. In a dark room




Figure 0.1.: Keith Emerson playing a modular Moog. The keyboard is connected to a big rack of modules, interconnected with each other with wires. Each module generates or processes the sound carried in analog form in the interconnecting wires, therefore having an exponential number of possibly different sounds by combining modules and settings.

only illuminated by a bright blue table, several performers placed tagged objects on the table. The table automatically arranged a connection graph among the objects based on their nature and relative distance and it displayed it but, more interestingly, the sound was evolving as this graph did. I just thought: wow, that must be fun, I want to play with that! — well, Do It Yourself.

That was the birth of Psychosynth. At the beginning it was just a toy. I did know nothing on how to process real-time audio, so I wrote many experiments. When I was bored

of testing the dynamic range of my speakers and ears with all sorts of sinusoids, triangle and square signals I started to read more and more source code of other audio Free Software and eventually started to write a digital modular synthesizer and play with Ogre3D.¹ By the summer of 2007 I had some very primitive synthesis code and also some user interface experiments². While I was an average C developer when I started my studies, I did not have any clue on C++. During the development of these initial experiments, I also had to learn about inheritance, what `virtual` means, etc. but of course the design was flawed all the time and I had to rewrite the code many times.

Something happened then at the beginning of my second year of university: I applied as a contestant to the Spanish Free Software Contest for University Students, and Psychosynth was the project I would develop. At that time I was already interested in experimental electronic music from the 90's, and late programming nights were accompanied by Autechre's fractal glitches and Aphex Twin and Squarepusher's spastic patterns³. At that time, this cruise in the most experimental side of electronic music and my ignorance in proper music making made me believe in scoreless generative music produced from evolving continuous signals, and that influenced the lack of proper synchronisation

1 Torus Knot Software Ltd. <http://www.ogre3d.com>

2 As this video can show, <http://blip.tv/file/325103>, the software was just a bunch of buttons with a 3D blue table like Reactable's and no sound.

3 If there is something I am grateful for during the early development of Psychosynth is the patience of my flatmates during those noisy programming nights. Psychosynth was long-time nicknamed "the ambulance siren sound maker" because it was only able to produce recursively modulated sinusoids.

mechanisms in Psychosynth. At some point, Shaker08⁴, a music producer and DJ from Malaga, developed some beat loops to distribute along with the software and helped in early showcase performances. He also tough me a lot on how music is traditionally made.

The project won a price in that Free Software contest and he got some reviews in blogs. it then became part of the GNU project — an attempt to assure its long-term development and that it would remain Free Software in the future.

After that, however, the development stalled a bit. I had big refactoring ideas that never got the motivation to be accomplished. In that seek for perfect code motivated by an increasing interest in programming language theory and functional and generic programming, I also became more and more conscious of the flaws of the code —i.e. the lack of synchronisation mechanisms, MIDI (Music Instrument Digital Interface) support, pluggable modules, patch persistence, etc. make any serious non-experimental attempt to make music with it very hard. During these last 2 years I have learnt much more on the music production workflow and terminology, a process parallel to a final step in opening my ears to current electronic music, specially Drum and Bass, Dubstep, and even some Minimal and Techno. During last summer I got a MIDI DJ controller that got me to better understand the limitations and possibilities of Psychosynth for music mixing and I became a casual contributor of the best Free Software DJ software: Mixxx [3]. Also, the people at ArtQuimia⁵, the music production school where Shaker08 was educated, became interested in the project and has of-

⁴ <http://soundcloud.com/shaker08>

⁵ <http://www.artquimia.net/>

ferred lending gear for testing and supervision, guidelines and insight from a musician point of view.

So here we are now, in the fall of 2010. Still quite ignorant in music making but pretending to be a “digital luthier” motivated by passion for music. And trying to turn all this personal game into a final master thesis project. Lets see how it goes. . .




Figure 0.2.: Artwork for the music of Kraftwerk, with robotic representations of the band members playing with electronic instruments. On the front we can see an old analog sequencer. As opposed to a synthesizer —this is, a virtual instrument, which is in charge of producing sound of whatever note you tell it to— a *sequencer* is an electronic score that can not produce sound by itself. The notes are programmed in the device —in an analog sequencer, by toggling switches and knobs— and it sends them at the appropriate time and duration to the electronic instruments (synthesizers) it is connected to. While current digital sequencers are very powerful, at that time they had important limitations that influenced Kraftwerk's robotic but charm sound. Kraftwerk was one of the first pop bands to produce its music entirely with electronic devices and is considered the father of modern electro and techno and has heavily influenced many other styles like house and hip-hop.

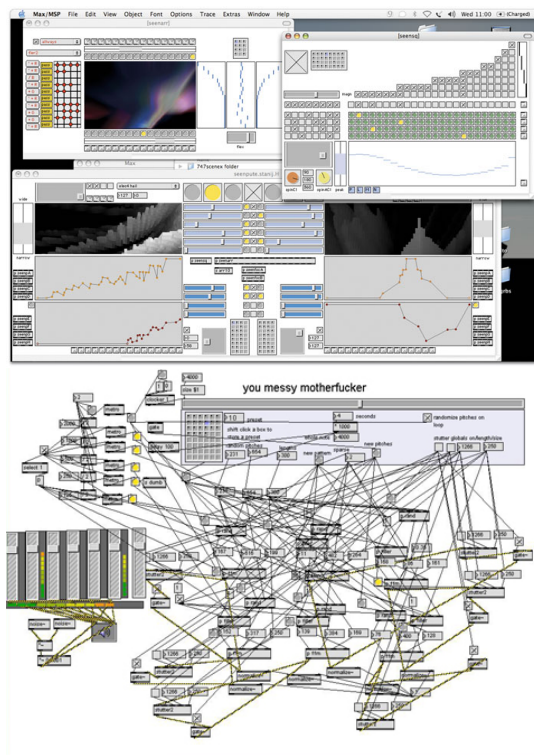


Figure 0.3.: The first picture is a Max/MSP claimed to be made by the electronic music duo Autechre[2]. The second is a Max/MSP patch made by Robbie Martin that *generatively* reverse-engineers Autechre’s song Powmod. Music is said to be *generative* when the whole composition itself is performed by pseudo-random algorithms programmed to produce well-sounding evolving non-repeating patterns. Autechre are known to have explored generative music. Max/MSP and its Free Software counterpart Pure Data are graphical data-flow programming software, which can be considered a low-level modular synthesizer, that is often used to write arbitrarily complex musical software and is specially interesting for generative music.

1

INTRODUCTION, DEFINITION AND GOALS

La utopía está en el horizonte. Camino dos pasos, ella se aleja dos pasos y el horizonte se corre diez pasos más allá. ¿Entonces para que sirve la utopía? Para eso, sirve para caminar.

EDUARDO GALEANO

1.1 PROBLEM DEFINITION

Our problem definition is well expressed in the title of this project: “a framework for modular, interactive and collaborative sound synthesis and live music performance”. Lets elaborate this by describing its parts.

1.1.1.1 *A modular synthesiser*

A modular synthesiser is one where the sound generation and manipulation process is described as a directional graph where each link represents the flow of sound signal from one node to another, and each node transforms or generates those signals.

A node with no input is often called a *generator*. A node with one input and one output is often called a *filter*. A node can have different *parameters* that in analog hardware can be set with knobs and potentiometers. Often, these parameter can be controlled with other optional input signals that are called *modulators*. A concrete interconnection of a set of modules is commonly referred as a *patch*. Note 1.1 lists some of the most common modules in analog modular synthesisers.

Note 1.1 (Standard modules in an analog modular synthesizer)

Software modular synthesizers have usually similar ones in their default module too, while they quite often use different terminology not related to this voltage based signal like in this case.

[The following text is extracted from the Wikipedia article on "Modular Synthesizer," as checked on December 12th 2010.]

VCO *Voltage Controlled Oscillator, which will output a pitched sound (frequency) in a simple waveform (most usually a square wave or a sawtooth wave, but also includes pulse, triangle and sine waves).*

NOISE SOURCE *A generator that supplies "hiss" sound similar to static, which can be used for explosions, cymbals,*

or randomly generated control signals. Common types of noise offered by modular synthesizers include white, pink, and low frequency noise.

VCF *Voltage Controlled Filter, which attenuates frequencies below (high-pass), above (low-pass) or both below and above (band-pass) a certain frequency. VCFs can also be configured to provide band-exclusion, whereby the high and low frequencies remain while the middle frequencies are removed.*

VCA *Voltage Controlled Amplifier, which varies the amplitude of a signal in response to a supplied control voltage.*

EG *Triggering an Envelope Generator produces a single, repeatable shaped voltage pulse. Often configured as ADSR (Attack, Decay, Sustain, Release) it provides the means to shape a recognizable sound from a raw waveform. This technique can be used to synthesize the natural decay of a piano, or the sharp attack of a trumpet. It can be triggered by a keyboard or by another module in the system. Usually it drives the output of a VCA or VCF, but the patchable structure of the synthesizer makes it possible to use the envelope generator to modulate other parameters such as the pitch or pulse width of the VCO. Simpler EGs (AD or AR) or more complex (DADSR—Delay, Attack, Decay, Sustain, Release) are sometimes available.*

LFO *A Low Frequency Oscillator is similar to a VCO but it usually operates below 20 Hz. It is generally used as a control voltage for another module. For example, modulating a VCO will create vibrato while modulating a VCA will create tremolo.*

RM *Ring modulator, two audio inputs are utilized to create sum and difference frequencies while suppressing the original signals. This gives the sound a “robotic” quality.*

MIXER *A module that combines multiple signals into one.*

S&H *Sample and hold, which takes a “sample” of the input voltage when a trigger pulse is received and “holds” it until a subsequent trigger pulse is applied. The source is often taken from a noise generator. Sequencer, which produces a sequence of notes, usually a music loop.*

SLEW LIMITER *Smooths off the peaks of voltages. This can be used to create glide or portamento between notes. Can also work as a primitive low-pass filter.*

CUSTOM CONTROL INPUTS *Because modular synthesizers have voltage-driven inputs, it is possible to connect almost any kind of control. Pitch can be varied by room temperature if you wish, or amplification varied by light level falling on a sensor.*

Analog modular synthesizers were invented in parallel 1968 by R. A. Moog Co. and Buchla in 1963 [4]. There, sound signal is more often represented by oscillating voltage levels running through wires—the links—and manipulated by analog signal processor modules. Usually these modules were arranged in big racks.

One of the biggest problems of modular synthesizers is their limited ability to cope with *polyphony*. We say that a synthesiser is polyphonic when several different notes can be played at the same time. The basic implementation tech-

nique for this is to have several copies of the synthesis logic—each is called a *voice*— and dispatch each keystroke to an unallocated voice (if present, otherwise, some note priority logic is to be implemented). In old analog modular synthesisers this was achieved by having multiple root oscillators, maybe 2 or 4, but this multiplied the complexity of connecting the synthesis graph. That is one of the main reasons why modular synthesis was gradually abandoned for analog devices, as a naturally sounding keyboard-controlled instrument should have a much higher grade of polyphony and should remain usable—note that the required polyphony level is higher than the maximum number of keys that we want to be able to press simultaneously, because a note remains playing after the key is released during a *decay time* while the sound softly fades out.

Nowadays, the increasing power of computers allows us to build modular synthesisers by software. Even further, we are no longer limited by wires and we can use arbitrary data types as processing and instantiate copies of the modules as we wish only constrained by our memory and computation power. On software, it is easier to achieve polyphony but it is still a non-trivial problem to make it efficiently.

1.1.2 *An interactive synthesiser*

Even though Keith Emerson virtuously manipulated the wires and knobs of his Moog in the middle of his performances, old-school modular synthesisers have the inconvenience that they are rather static. It is very hard to control all the parameters during the performance. Changing the topology of the synthesis network is almost impossible, and in

most systems it causes clicks and other inconvenient noises as a result of abruptly connecting and disconnecting the wires —whether software or analog.

We should design our engine with care so no noise is generated as a result of manipulating the system live. But we should also provide other means to enable easier manipulation of the topology.

1.1.2.1 *Dynamic patching*

The *dynamic patching* [5] technique was first introduced in the ReacTable project and offers means to automatically interconnect the modules of a software modular synthesiser. When present, modules are laid out in a bi-dimensional space and each output automatically connects to the nearest available input of its same kind. The sink node that leads the output to the speakers is situated in the centre and the synthesis graph grows as a radial tree around it, as shown in figure 1.1. By simply moving one module from one position to another the user can radically change the topology of the network without doing a lot of plugging and unplugging of wires. A clever disposition of the modules on the space can help the artist to achieve new innovative ways of producing variations in his music.

1.1.2.2 *Touchable interfaces*

An specific problem for music software is that its interface is highly limited by the keyboard and mouse as interface. While a person has 20 fingers¹ she is limited to manipulate only one parameter at a time with the mouse.

¹ She must be very skilled to use all of them at the same time!




Figure 1.1.: A screenshot of Psychosynth 0.1.1 using dynamic patching to connect a complex graph.

There is an increasing availability of touchable interfaces, either specifically designed for music performance like the Lemur (figure 1.2) or general purpose ones like the iPad. There, one is no longer constrained by the single-click paradigm and not only can she manipulate various parameters at a time, different multi-tap gestures can expand the possibilities of a spatially limited user interface as several functions can be reached in the same points.

1.1.2.3 *Instrument simulating controllers*

While touchable interfaces might be better for manipulating the continuous parameters of the synthesis and dynamically manipulate a sequencer, many musicians would prefer more traditional interfaces to interpret part of their work in the “old-fashioned” way of playing an instrument [6]. There exists many kinds of electronic devices that simulate the feel of a keyboard or a drum-kit (figure 1.3) but instead




Figure 1.2.: The JazzMutant's Lemur touchable music interface. All the on-screen controls can be configured with an UI designer program, and later mapped to any program via OSC messages.

of producing any sound they send MIDI [7] messages to a synthesiser that reproduces the notes.

Our software should be able to interpret MIDI messages such that it can be controlled with such a device, making it more natural and creative for many musicians.

1.1.3 *A collaborative environment*

Since the beginning of times, music performance have been a collaborative act, with each performer contributing to the global rhythm and harmony by playing one instrument. Psychosynth should serve the purpose of producing arbitrarily complex music by itself as modules that implement not only synthesis, but also sampling and sequencing are added.

This integrated environment should be able to be manipulated by several people at the same time to allow a



Figure 1.3.: A conga-like MIDI controller. When pressing different parts of its surface it will emit different MIDI note messages, that a synthesiser or sampler could use to either simulate a real conga or to produce any other sound.

collaborative experience. After all, it would be very hard for one person to control in a live performance all the subtleties of the sound generation by herself.

1.1.3.1 *Tangible interfaces*

One approach to achieve this is by using a user interface that is big enough to accommodate several performers around it. A tangible interface is one where the different elements are represented by physical objects that one can touch and move in physical space. The ReacTable implements such an interface where the modules of its synthesis engine are plastic blocks that are laid out over a round table that provides visual feedback thanks to a projector under it. As shown in figure 1.4, with such an intuitive and physically unconstrained interface several people can stand around the device and manipulate it.



Figure 1.4.: An example of the Reactable being used collaboratively by several people.

1.1.3.2 *Networking*

Networking support even further releases a device from the space constraint by allowing several instances of the software intercommunicate over a computer network —i.e. IP. At some point, latency problems can still be a drawback for this technique, but it can be useful for some kinds of collaboration that do not require perfect timing. When playing in a local range, this becomes perfectly valid even under high latency requirements.

This has long time been a main goal in Psychosynth. When running in collaborative mode, all the clients that connect to a server share the same synthesis graph and whenever an object is moved, added or deleted, or a parameter is changed, all are notified such that they keep the same view of the scene. Figure 1.5 shows an example of this feature being used live.



Figure 1.5.: An example of Psychosynth 0.1.4 being used collaboratively over the network. The picture was taken in a showcase performance held by Shakero8 and the author in the Open Source World Conference in 2008.

1.1.4 *A framework*

Music making and audio processing software and interfaces are evolving very fast. Abstracting the common core functionality needed by developers in a layered and abstracted programming API and development of a Free Software [8] framework is crucial to enable further focused research on the different areas previously discussed. Results of that research can be later integrated in the framework as they are stabilised.

Such a framework should enable the development of modular audio applications with an abstracted interfacing mechanism, probably following a Model-View-Controller paradigm, general enough to be able to support all the previously described qualities. If this task is properly accomplished, the framework could become the basis of a wide range of unexpected future Free Software applications.

1.2 OBJECTIVES

Taking all that into account, we should next define the concrete objectives for our present project. We should note that we depart from the basis of the current state of the GNU Psychosynth project—as of its version 0.1.7—and we assume that as previous work, not our current target.

1.2.1 *Main objective*

Objective 1 *Re-focus the GNU Psychosynth project as a development framework for the development of professional-quality modular, interactive and collaborative synthesisers.*

In the long-term we would like GNU Psychosynth to include innovative user interfaces, and some might even be developed as a side effect of this project—or we might just update the older one to rely on the new framework. However, that is not the purpose of this project, instead we will concentrate on the development of the underlying core architecture and implementation of its API.

This is so mainly because of time constraints. Also, if we were to miss some features in the core in order to allocate more time for the user interface, we have to take into account that this will probably hard to fix afterwards when a lot of code depends on the broken design. Also, user interface development is easier to do in a non-disciplined, voluntary and heterogeneous team. If we achieve a nice framework now we can still develop the user interfaces later with the help of other the people collaborating on the Internet; but, as these two years of stalled development have shown, it

is hard without the pressure of an external deadline and a project plan to invest a lot of time in rewriting the “invisible” but crucial parts of the system.

1.2.2 Preconditional objectives

There are two objectives of this project that can also be considered as a precondition for the success of our main goal. These are:

Objective 2 *Collaborate with professional musicians to get a defined understanding of the meaning of “professional quality” and their real needs.*

The students participating in this project have an amateur knowledge of music production. It is important to communicate and allow supervision by professional musicians and experts in digital music to assure the suitability of the software for use in a professional environment.

We are working in collaboration with the ArtQuimia Music Production School², which has long time been educating successful producers and is currently participating in the European Cross-Step project³, and his director David García, a musician with professional experience in the industry as music composer and sound designer for video games.

Me consta que David ha trabajado en la industria de videojuegos sintetizando sonidos pero debería preguntarle a ver si hay algo más destacable en su currículum.

² <http://www.artquimia.net/>

³ <http://www.cross-step.org/node/4>

Objective 3 *Research and apply the latest techniques in modular design and implementation and explore the boundaries of the underlying implementation devices.*

The success of an framework relies on the proper decomposition of its features and its extensibility. Moreover, the authors of this project have a special fascination for programming languages, design patterns and software modularity. Even more, there is an active research community questioning and re-developing the modularisation and abstraction techniques of the underlying programming language C++, a fact that is more true as we approach the final resolution of the standardisation committee on the new C++0x standard. All this suggests that research and application of the state-of-the-art and even development of our new design and coding patterns will be one of the leading objectives during the project and play a leading role in its overall success.

1.2.3 *Partial objectives*

A more concrete subdivision of our main goal should be given. Note that these are not yet the detailed requirements, but an overall initial objectives vaguely elicited from the problem definition, the previous experience with the GNU Psychosynth software and our personal interests.

Objective 4 *Improve the framework to be able to load modules dynamically from plug-ins, satisfying our own API and/or third-party standards.*

This is a common feature in most industry standard applications and ours should support it. Many layers of the

framework, specially those related to the dynamic patching, will require vast modifications to enable customisation to understand third-party modules.

Objective 5 *Improve the framework to be able to communicate with music controllers via MIDI or other industry standards.*

While not explicit in this wording, this adds the requirement for *polyphony* as in most cases such feature would be useless without it.

Objective 6 *Add basic synchronization and sequencing support to the framework.*

If we want to understand the software as a full live performance environment and not a bare synthesiser, this currently lacking feature is fundamental.

Objective 7 *Include common modular audio system utilities into the framework. Some of the most important being patch hierarchy and persistence.*

1.3 BACKGROUND AND STATE OF THE ART

1.3.1 Modular synthesis

The history of modular audio generation starts with the Moog analog synthesiser in 1967 [4]. Since then, a wide variety of analog modular synthesisers have been developed commercially, but retaining some limitations as described in the introduction in section 1.1.1.

Modular synthesis became more interesting with the up-rising of computer based electronic music. One of the most

important examples in this development is Max/MSP [9]. This is a dataflow based visual programming environment mainly targeted at music applications. In such software one can add boxes where one types the name of the function that it should perform. When the name has been written some connection plugs appear on its corners depending on the module name, and one can draw lines connecting those plugs. Figure 0.3 showed an example of its functioning. The author of Max/MSP later developed a Free Software counterpart called Pure Data[10] that also has video processing and synthesis support.

A continuación hablo de aplicaciones comerciales que no tienen (o no he encontrado) artículos académicos. Estoy evitando citar los manuales de usuario o simplemente enlazar a sus webs comerciales, pero no sé que es lo más ortodoxo hacer en estos casos. En algunos casos más adelante y cuando lo he considerado relevante estoy poniendo enlaces al pie de página cuando no existe un paper académico. Cualquier recomendación al respecto es bienvenida.

Since then, many user-oriented commercial modular synthesisers have been developed. One of the most famous ones is Native Instrument's Reaktor (figure 1.6). In its version 5 it included the new *core-cell* technology, which allows the visual and modular design of the lower level parts of the DSP (Digital Signal Processing) programming, which transparently compiled to efficient machine code. Later, Plogue's Bidule is gaining special recognition for its simpler interface. A new interesting software is XT Software's EnergyXT, a DAW (Digital Audio Workstation) that has a "modular view" where one can route through wires all the MIDI and audio signals that flow behind the traditional sequencing view.

Finally, Sensomusic's Usine is remarkable for introducing a highly customisable multitouch graphical interface on top of a modular synthesis environment.

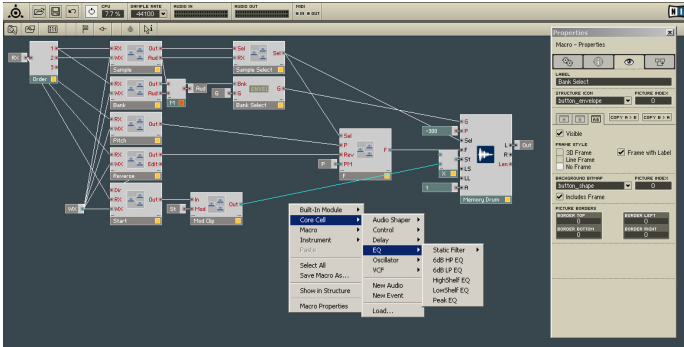


Figure 1.6.: Native Instrument's Reaktor editing a *core cell* based patch.

On the Free Software side few modular synthesisers exist. Also Modular Synth⁴ is one of the most popular. Ingen⁵ is a more modern one whose code we praise for its quality. Most of these software still lack some of the features of their privative counterparts, with the ability to create customised user interfaces being the most relevant. Still, Free Software offers a very interesting modular approach to the sound system management that has a similar rival only in latest OSX versions: Jack[11]. Jack is an audio server providing zero-latency interprocess communication that acts as a patch-bay among any supporting audio applications. The user can route the output of one program to the input of any other program, or the soundcard sink, or whatever exposes a Jack port. It can be used to route MIDI signals too and

⁴ <http://alsamodular.sourceforge.net/>

⁵ <http://drobilla.net/blog/software/ingen/>

recent versions include very nice features, like collaborative routing over the network for distributed performances [12].

1.3.2 *Touchable and tangible interfaces*

In the last decade, the development of touchable and tangible user interfaces have been of rising interest. An interesting but not very updated survey that gives some taxonomical background and analyses a wide variety of products can be found here [13]. One of the first attempts in using them for improved interaction in musical software is the Audio Pad[14], where the user places and moves tagged pucks on a bidimensional surface. This is an example of an interface with active tangibles, because the pucks have an RF transmitter to allow their recognition. The Jam-o-Drum, on the other hand, offered a percussive collaborative approach where performers sit on the corner of an interactive table [15]. Many videogames and other kinds of application have been developed on top of the Jam-o-Drum hardware too. Since then, a huge number of different table based interfaces have been made, an example listing can be found here⁶.

Maybe the most interesting example, which inspired the whole development of GNU Psychosynth, is the Reactable project[1]. Its user interface is based on modular synthesis and uses the dynamic patching technique for easily manipulating the patch topology. It uses innovative digital image processing techniques to detect the position and rotation of passive tangibles [16]. In the Reactable system a translucent round surface hides a projector and a high resolution camera underneath, as shown in figure 1.7. Finger-tips and the

⁶ <http://mtg.upf.edu/reactable/related.htm>

specially designed tags called *fiducials* that are placed on the table surface are captured by the camera and recognised by the ReactIVision system on a computer. This system sends OSC[17] (Open Sound Control) based messages to a synthesiser following the TUIO (Tangible User Interface OSC) protocol [18]. The literature describing the initial prototypes used Pure Data as the underlying implementation language for this synthesis engine, but conversations with the authors of the Reactable suggest that the current commercial implementation is written in C++. This synthesis engine is connected to a visual engine that generates visual feedback and sends it through the projector. The picture is reversed so it can be correctly visualised through the translucent surface.

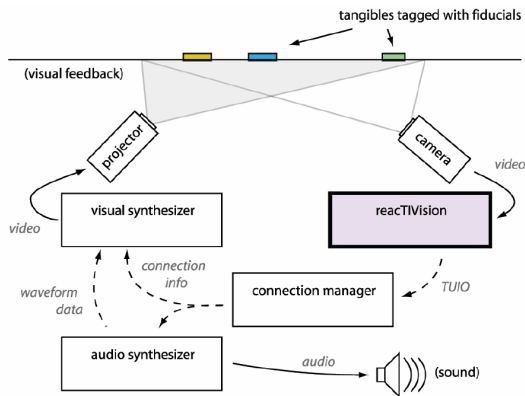


Figure 1.7.: A typical Reactable setup. Source: The Reactable project.

The Reactable has been very successful and a very interesting fact is that the computer vision component of the system is Free Software, so it could be integrated with GNU Psychosynth in the future.

Some other remarkable tangible and highly interactive user interfaces that were released around 2007 too are the JazzMutant's Lemur and the Iwai-Yamaha's Tenori-On [19]. The former offers a fully customisable OSC based multi-touch interface, where one can design interfaces with sliders, virtual-knobs, X-Y pads and all kinds of multi-touch controls that are mapped to OSC messages that can be interpreted by the audio engine of choice. The latter is a 16x16 matrix of LED illuminated buttons that can be configured in different modes that provide innovative sequencing mechanisms (figure 1.8).



Figure 1.8.: Electronic music artist Four Tet performing with a Tenori-On

Nowadays, multitouch interfaces are the main trend, specially after the explosion of the tablet market, but big music software companies still have to catch up the latest hardware developments. Specially interesting for the GNU Psychosynth projects is the Indamixxx 2 tablet, which is based on the Meego OS⁷ and oriented towards music production

⁷ Nokia and Intel's joint effort to provide a GNU/Linux based operating system for embedded devices, mobile phones, netbooks and tablets.

and live performance — we should definitely keep an eye on it and develop a multi-touch enabled user interface for it.

1.3.3 *Audio synthesis libraries and frameworks*

One of the first synthesis libraries that were evaluated before the development of GNU Psychosynth started is the STK (Synthesis ToolKit) [20] but it lacks proper dynamic signal routing mechanisms and some details, like the sample format being hard-coded to float, seem too constrained.

A popular sound synthesis and music composition DSL (Domain Specific Language) is CSound, which later was extended with a C and C++ programming API[21]. Many DSP oriented DSLs have been developed, but they are not general enough to support the development of the fully featured applications that we wish on top of GNU Psychosynth. Still, some of them are worth mentioning. Notable examples are SoundCollider [22], which was for a long time the reference sound synthesis DSL; Chuck [23], that adds concurrency and time control in an elegant and boilerplate-less fashion, and the newer Faust [24], which is based in a functional and data-flow paradigm and interfaces easily due to its compilation to C++.

While we are focusing on C and C++, many languages have music related libraries. Impromptu⁸ is a Scheme environment for live-coding, this is, live music performances where the programming is done in front of the audience while the programmer writes and mutates the code⁹. It sup-

⁸ <http://impromptu.moso.com.au>

⁹ The video *Algorithms are thoughts, chainsaws are tools*, by Stephen Ramsay is a great introduction to this technique.

ports video manipulation too, but sadly is available for OSX only. Its heavy modularity and dynamism is what make Lisp dialects so interesting for live-coding and music composition. Common Music[25] was started in 1989 and it is a highly featured Common Lisp framework that, among other things, provides a very elegant embedded score notation.

Maybe the most similar project to GNU Psychosynth in its approach is CLAM (C++ Library for Audio and Music) and award winning library based on modular synthesis[26]. Its design will be carefully taken into account during the redesign of few of Psychosynth core components. Still it does not precisely fit our needs as it is too oriented toward research simulations and does not support polyphony.

<http://createdigitalmusic.com/2010/07/thought-and-performance-live-coding-music-explained-to-anyone-really/>

2

ANALYSIS AND PLANNING

As I see it, criticism is the prime *duty* of the scientist and of anyone who wants to advance knowledge. Seeing new problems and having new ideas, on the other hand, are *not* one's duty: originality is, rather, a gift of the gods.

1982 Preface to *Quantum theory and the schism in physics*

KARL POPPER

2.1 REQUIREMENT MODELLING

In the following section we discuss the requirements that we want to impose on the resulting framework. Note that functional requirements are hard to express for an API, but a vague description of the desired qualities and features is still possible will guide or development process.

Quite often, we will focus on the final feature that should be implementable through the framework. Also, note that the framework already had many features at the beginning of the project, some of which we will discuss in section ??.

We will try to avoid describing the requirements related to those ready facilities in this section, as the only purpose of those is to aid the actual development of this thesis project. Still, some already satisfied requirements might be made explicit when some other related requirement follows, or because of their relevance to the user or just for the global consistency of the text.

As we suggested in our objective 2, these requirements have been elicited with the assistance of the ArtQuimia Music Production School in order to ensure the suitability of the software in a productive usage context.

2.1.1 *Functional requirements*

2.1.1.1 *Basic audio support*

Requirement 1 *The library must provide data-structures and basic algorithms for manipulating audio signals.*

Requirement 2 *The library must provide means to output and input audio data to and from files in various formats, including, at least: Microsoft WAV, Apple AU and OGG Vorbis.*

Requirement 3 *The library must provide means to output and input audio data to and from sound-card devices in an asyn-*

chronous real-time fashion, supporting, at least, the following audio systems: ALSA¹, OSS², Jack³.

2.1.1.2 Node graph support

Requirement 4 *The library must include means for producing the audio as the result of executing a dataflow-graph.*

Requirement 5 *The library user must be able to define his own processing nodes.*

Requirement 6 *Each node should have an arbitrary number of named signal and control input and output ports. The difference between signal and control ports are that the later are sampled at much less frequency.*

De hecho, la naturaleza real de las señales de control aún no la tengo clara al 100%. Las señales de control son lo que en la versión actual de GNU/Psychosynth se llaman *parámetros*. Sin embargo, el tema no es del todo trivial y tras ver como funciona Reaktor y pensar mucho en como reinventar el MVC del Psychosynth la situación está cada vez más confusa, pero supongo que la concreción se puede dejar para el rediseño de esta capa del software en la segunda iteración.

Requirement 7 *Both signal and control ports should be able to process information in arbitrary datatypes. Signal ports may have practical limitations as for the real-time constraints is concerned.*

¹ Advanced Linux Sound Architecture: <http://www.alsa-project.org/>
² Open Sound System: <http://www.opensound.com/>
³ <http://jackaudio.org/>

Note 2.1 (Realtime constraints)

All the processing done inside the nodes should satisfy soft real-time constraints. This is so because in order to produce sound with low latency (see requirement ??) the sound processing should be done in blocks (buffers) as small as possible that are delivered to the sound-card as soon as possible. If the deadline is not met, a disturbing “click” sound will be heard because of the zeroes assumed by the sound-card during the period that it did not have any audio to output. For example, for a 44100 Hz sampling rate and a 128 frames block size, it should take less than

$$\frac{1\text{s}}{44100 \text{ frames}} \cdot 128 \frac{\text{frames}}{\text{block}} = 2.90 \frac{\text{ms}}{\text{block}}$$

to process and deliver an audio data block.

In practice, this means that the processing algorithms should take an amount of time proportional to the amount of data to process —i.e. they are $O(n)$. This in turn disallows writing and reading files and most operations that cause non deterministic hardware operations or force context switching (mutexes might be unavoidable but should be used scarcely). Also, allocation of memory on the heap should be avoided as it has a non-predictable and potentially non-linear time consumption. Of course, the framework should provide hooks to do those forbidden operations outside the audio processing thread, but we consider that a design issue that should be addressed later.

Requirement 8 *Each output port must be connectable to arbitrary number of input ports. Each input port must be connectable to one output port.*

Requirement 9 *Ports may be defined statically —i.e. at compile time— or dynamically —i.e. at runtime.*

Requirement 10 *The system must allow the hierarchical composition of nodes, with special input and output nodes that are exposed as ports in the parent level.*

Requirement 11 *Nodes can be either monophonic or polyphonic. A polyphonic node internally has a number of copies of its processing state, called voices, that are dispatched accordingly through trigger signals.*

Note 2.2 (On the scope of polyphony)

We will avoid specifying here details on how polyphony works that are still quite important as a usability concern. For example, should monophonic ports be connectable to and from polyphonic ports? Should there be polyphonic and monophonic nodes in the same patch hierarchy level? How are the voices dispatched and mixed? Because answering this issue highly affects performance and implementability trade-offs, these issues are left open until the design stage of these components.

2.1.1.3 Dynamic loading of nodes

Requirement 12 *The system must be able to dynamically load modules developed with standard interfaces, at least using the*

LADSPA standard⁴, but LV2⁵ and VST⁶ are proposed too in this order of priority. Note that in some cases this implies the automatic fixing of the semantic impedance among the interfaces of Psychosynth nodes and the third-parties’.

Requirement 13 *The system should be able to dynamically load modules developed using the same native interface exposed to library users to define their own modules.*

2.1.1.4 Dynamic patching

Requirement 14 *The system should optionally release the library user from arranging the interconnections among nodes by using dynamic patching. When using dynamic patching an output port is connected to one input port. The input port is the closest (in the sense of euclidean distance, we assume that the modules are laid out in the space) input port of the same type not belonging to the same node that is free —i.e. is not connected already to a closer node.*

When combining dynamic-patching with dynamically loaded modules using standard interfaces like LADSPA, one might need extra information to correctly perform the dynamic patching.

Requirement 15 *The system should be able to locate and automatically load if present special description files containing the required information for some dynamically loaded modules to work correctly. This file should specify which ports are eligible and to which ports they are connectable to.*

4 Linux Audio Developer’s Simple Plugin: <http://www.ladspa.org/>

5 LADSPA Version 2: <http://lv2plug.in>

6 Steinberg’s Virtual Studio Technology: <http://ygrabit.steinberg.de/>

While this may change due to design considerations, we suggest specifying a set of tags for each port, with the following connectability rule for port A and B : $connectable(A, B) \Rightarrow tags(A) \cap tags(B) \neq \emptyset$.

2.1.1.5 *View independence*

The following requirements are here to suggest an observable interface for the synthesis model satisfying a model-view-controller architecture. This is one of the key concepts in achieving a network-based shared environment that is transparent to a wide range of graphical interface experiments or external MIDI controls.

Requirement 16 *The system must enable the development of graphical interfaces —or any other instance of the more abstract concept of view— that can coexist with each-other.*

2.1.1.6 *Synchronisation and sequencing*

Requirement 17 *The system should include a notion of tempo such that external imprecise events can be quantised and synchronised to.*

Requirement 18 *Parameters of various node kinds, specially time related ones, should be controllable as a factor of the current tempo.*

I.e. the frequency of an LFO should be fixable such that the wave length is a factor of the tempo, and the phase should be synchronisable to beat.

Requirement 19 *The system should be able to synchronise to the MIDI clock.*

2.1.1.7 Collaborative support

Requirement 20 *The system must be able to receive, process and use as control source MIDI events coming from other software or external hardware devices.*

Requirement 21 *The system must be able to receive, process and use as control source OSC events coming from other software, computers or external hardware devices.*

Requirement 22 *The system must be able to use a specially crafted OSC based protocol to enable the collaborative manipulation of the same node graph among different computers connected through the network.*

Note 2.3 (Relaxing the “shared-environment” requirement)

Requirement 22 is currently implemented by broadcasting all events among all the participants in a shared session. In presence of audio input nodes, dynamically loaded modules and other sources of non-determinism —this is, sound that might not depend only on the sequencing of certain events— this gets harder to implement. We do have some solutions in mind, like placeholder nodes that stream the local non-deterministic signals through the network too, but it might be too hard to implement in the context of this master thesis project and only a proof-of-concept implementation will be required.

2.1.1.8 *Persistence*

Requirement 23 *The system must be able to store and load the graph layout and control values the node graphs. Sub-graphs in the graph hierarchy should be storable and loadable individually.*

2.1.1.9 *Optional functionality*

Requirements in this section are not such, but instead they are bare ideas that would be nice to have but are considered too time consuming, hard or not urgent enough to be considered a measure of the project success.

Requirement 24 *The highest level part of the API should have a Python —or any other dynamic language of choice— binding for the rapid prototyping of audio applications and user interfaces on top of it.*

2.1.2 *Non-functional requirements*2.1.2.1 *Free Software*

Requirement 25 *Unless constrained by very specific hardware, the system should not add any non-Free Software dependency — i.e, it must be able to compile and run without executing any privative software bit.*

2.1.2.2 *Professional quality sound*

Requirement 26 *The software must be able to work at different sampling rates up to 96000 Hz.*

Requirement 27 *The software should be able to use arbitrary precision samples. Up to 64 bit floating point samples are required.*

2.1.2.3 Performance

Requirement 28 (Latency) *The software should be able to work with a block size as low as possible down to 64 frames, as long as the underlying hardware permits it.*

En realidad este requisito es un poco ambigüo, pero no se muy bien cómo especificar la posibilidad e usar tamaños de buffer bajos, ya que las condiciones en las que es medible esto dependen de la potencia del procesador y, en un sinte modular que permite redes arbitrariamente complejas, del tamaño del sinte. ¿Alguna sugerencia?

2.2 OPEN, COMMUNITY ORIENTED DEVELOPMENT

The developers of this software are strong proponents of Free Software development. This is specially relevant in an academic and public environment like ours. Therefore, the software not only is distributed under a Free Software license, it also follows an open community development model, where everyone can read, use and modify the source code as it is developed and there are means for online communication promoting development among volunteering distributed peers.

Previous versions of the software are available on the Internet for download and it has an official web page: <http://www.psychosynth.com>

2.2.1 *Software License*

The software is licensed under the GPL license version 3, offering strong copyleft requirements — i.e. derived works and software linking against the library must be distributed under the same terms. The full description of the license is included in the appendix B.

While the GPL₃ is often misunderstood as inadequate for a library, that is not true in the context of libraries that provide unique features, as it motivates the release of third-party software that is attracted by these as Free Software too [27]. This is not only personal belief, it is also the official guideline in the GNU project.

2.2.2 *The GNU project*

Since October 2008, Psychosynth is part of the GNU project. GNU was started in 1984 by Richard Stallman with the long term goal of providing a fully free —as in speech— operating system [8]. Under the umbrella of GNU, Psychosynth gets access to a broader community, technical support and it is a recognition of its attachment to the Free Software ideals.

2.2.3 *Software forge and source code repository*

A *software forge* is an online service targeted at aiding the development of Free Software. It offers a series of tools to aid the community participation and distributed development, such as *bug trackers*, *support trackers* and *mailing lists*. One of

the most important features is the *revision control system* that serves as source code repository.

GNU Psychosynth is hosted in the Savannah software forge, and its project page can be accessed here:

<http://savannah.gnu.org/projects/psychosynth>

The project is using GNU Bazaar as distributed revision control system. One can fetch a copy of the latest version of main development branch by executing the command

```
bzr branch http://bzr.sv.gnu.org/psychosynth/trunk
```

2.2.4 *Communication mechanisms*

Fluent distributed communication is essential for the advancement of a free project. For this purpose we offer the following tools.

2.2.4.1 *A blog*

The blog serves as an informal and easy way of getting the latest news on the development. It is most of the time technically oriented and can serve as source of motivation for external people to contribute to the project and as a summary of the current status of the project. It can be accessed through:

<http://planet.gnu.org/psychosynth/>

2.2.4.2 *Mailing lists*

Mailing lists are multi-directional broadcast based communication means and the main spot for discussion of develop-

ment (from the developer point of view) and getting news or asking for help (from the spare user point of view).

GNU Psychosynth has two mailing lists.

- Users mailing list:
<http://lists.gnu.org/mailman/listinfo/psychosynth-user>
- Developers mailing list:
<http://lists.gnu.org/mailman/listinfo/psychosynth-devel>

Because being registered in many mailing lists can cause management issues to some users, the Gmane project⁷ offers a newsgroup gateway that can be used by Free Software projects to allow participation in their mailing lists with Usenet capable software. Psychosynth mailing lists are registered there and can thus be accessed through:

- `gmane.comp.gnu.psychosynth.user` for the users mailing list.
- `gmane.comp.gnu.psychosynth.devel` for the developers mailing list.

2.3 DEVELOPMENT ENVIRONMENT

The development environment is very important for the project success. This section should clarify our choices and explain the rationale behind such decisions.

⁷ <http://www.gmane.org/>

2.3.1 *Programming language*

Psychosynth is developed using C++. This decision is based on the following facts:

- It is a stack based language with fine grained control over memory management. As we introduced in note 2.1, this is crucial during the development of live audio processing software.
- It is a mature language with widespread tool support and a very good Free Software compiler: GCC.
- Apart from its low-level mechanisms it has powerful means for abstraction, most of which are designed to pay zero cost.
- It is multi-paradigm, and as such it can easily adapt to the natural way of expressing the concepts of a heterogeneous system as this, where we want to go to from low level DSP programming to high level interface design.
- It is compatible with C, which gives us direct access to the widest range of libraries available. Most audio-processing libraries are written in C and thus we can save a lot of time in implementing mathematically complex algorithms.

Of course, it also has its flaws, like unnecessary complexity and unsafety in some corners, most of which are justified by its backward compatibility to C and evolutionary design. Some of this flaws are nonetheless going to be solved in the next standardisation of the language, to be released in 2011 [28].

Aquí estoy citando un blog — ¿Es esto correcto o mejor lo pongo como nota al pie?

Compilers are starting to support it and we are very interested in exploring the benefits and new programming patterns and design benefits that it can provide. Because Psychosynth is an ongoing and forefront project we do not fear portability — and after all, GCC is very portable! — and as such we are going to use the facilities in C++0x as soon as they are supported by the latest version of GCC included in the Debian Sid.

2.3.2 *Operating System*

The project is mainly targeted at GNU/Linux, which is the most widespread free operating system, therefore, compliance with it is the highest priority. That is also the operating system of choice of the authors of this project so it feels like a natural environment and there is no extra effort needed.

Still, we will try to comply with POSIX such that porting to other Unix operating systems is easy. Sadly, there is no universal high performance and fully featured cross-platform audio output engine and thus that is an important portability boundary. In the future, maybe with some financial aid, we might be able to port the software to OSX, which is near to be the most used operating among musicians [6]⁸.

⁸ The cited survey dates back to 2006. Given the recent rise in popularity of Apple products, we speculate that OSX might be even more popular than Windows among musicians nowadays.

2.3.3 *Third party libraries*

In this section we give an overview of the external libraries used in the software. Note that they have to be chosen in compliance with requirement 25.

2.3.3.1 *Boost*

Boost⁹ is a set of libraries for C++ that are peer-reviewed and specially crafted to integrate well with the paradigms and abstraction techniques of the standard library. Many of its modules are, actually, going to be part of the standard library in the future C++0x standard.

We use few of the Boost facilities, some of them including `boost::filesystem`, `boost::thread`¹⁰, `boost::mpl`, `boost::test`¹¹ among others.

It is extremely portable and licensed under the permissive Boost Software License.

2.3.3.2 *Libxml2*

We use Libxml2¹² for parsing the XML configuration files. It is very portable and licensed under the permissive MIT License.

⁹ <http://www.boost.org>

¹⁰ We plan to replace this with the new threading facilities in the standard as soon as they are implemented in GCC.

¹¹ A wonderful unit testing framework.

¹² <http://xmlsoft.org/>

2.3.3.3 *LibSndFile*

We use LibSndFile¹³ for loading different uncompressed sound formats. Note that from version 1.0.18 it also supports OGG and FLAC formats, and thus we plan to use this instead of LibVorbis in the future. It is very portable and is licensed under the LGPL 2 and 3.

2.3.3.4 *LibVorbis*

We use LibVorbis¹⁴ for reading OGG files. It is released under the BSD license and is very portable.

2.3.3.5 *SoundTouch*

We use SoundTouch¹⁵ for time-stretching, this is, changing the pitch and tempo of a song independently. It is licensed as LGPL and works on mayor operating systems.

Some people claim to obtain better results in performance and sound quality with the Rubber Band library^{16 17} and we will probably replace SoundTouch by this one in the future.

¹³ <http://www.mega-nerd.com/libsndfile/>

¹⁴ <http://xiph.org/vorbis/>

¹⁵ <http://www.surina.net/soundtouch/>

¹⁶ <http://www.breakfastquay.com/rubberband/>

¹⁷ The DJ software Mixxx, where sound stretching quality is very relevant, is moving towards RubberBand and their developers support the above claims: <http://www.mail-archive.com/mixxx-devel@lists.sourceforge.net/msg03103.html>

Also, Ardour later moved tater moved to this library: <http://www.ardour.org/node/1455>

2.3.3.6 *ALSA, OSS and Jack*

In accordance to requirement 3 we use ALSA, OSS and Jack respectively. Licenses are LGPL2+ for all of them. ALSA is Linux only, Jack works on GNU/Linux and OSX and OSS works on Linux and some other POSIX operating systems.

2.3.3.7 *LibLO*

LibLO¹⁸ is used for OSC support. It is licensed as LGPL2+ and is POSIX compliant.

2.3.3.8 *The user interface libraries*

The user interface included at the beginning of the project is based on Ogre3D¹⁹, using OIS²⁰ (Open Input System) for keyboard and mouse input and CEGUI²¹ as widget toolkit.

During the development of the previous Psychosynth versions CEGUI has proven to be extremely painful with an overengineered API and bugfull implementation. Also, the 3D interface, while being fancy it is confusing for some users and did not offer anything new to the experienced musician, as some blog reviews showed. It contains some interesting concepts — zoomability being the most important — but it is too distracting after-all. Thus, while we are going to keep this interface during the development of this project, we want to later rebuild the GUI using Qt, which is multi-touch

¹⁸ <http://liblo.sourceforge.net/>

¹⁹ <http://www.ogre3d.org/>

²⁰ <http://www.ogre3d.org/>

²¹ <http://www.cegui.org.uk/>

enabled and could open us the door to the yet-to-come wide range of Meego based tablets.

2.3.3.9 *Build system*

We use GNU Autotools²² as the build system for the project. It is extremely portable and the de-facto standard among Free Software, even though some interesting alternatives are emerging. Nonetheless, Autotools are suggested in the GNU Coding Standards[29] that we shall follow during our development due to our affiliation to GNU.

2.3.3.10 *Documentation system*

Because we are developing a framework, it is specially important that the public interfaces are properly documented. We are going to use Doxygen²³ to embed the API documentation in code comments. A reference manual generated by Doxygen should be attached to this document.

2.4 ARCHITECTURE AND CURRENT STATUS

The Psychosynth project was started in 2007 and since there has been some relevant developments. The screenshot in figure 2.1 gives a broad perspective of the current features of the project. For a more detailed description from the user point of view, please refer to the user manual that can be

²² We use mainly Autoconf (<http://www.gnu.org/software/autoconf/>), Automake (<http://www.gnu.org/software/automake/>), and Libtool (<http://www.gnu.org/software/libtool/>).

²³ www.doxygen.org

checked online on the project webpage²⁴, where there are also demonstration videos and other multimedia material. The article [30] can serve as an introduction to the project too.




Figure 2.1.: A screenshot of GNU Psychosynth 0.1.4. On the bottom we can see some buttons for popping up different menu windows, such as for settings, recording the sound into a file or connecting several instances over the network. The window on the top of the screen is used to add elements to the processing graph. The smaller *properties* windows contains a listing of all the actual parameters of a node and allows us to give numeric values to them. All around the screen, sinusoids, step sequencers, file samplers and echo filters are interconnected as represented by the 3D blocks.

²⁴ <http://psychosynth.com/index.php/Documentation>

2.4.1 *A critique on the term framework*

The term *framework* is used many times in this and other projects and is becoming a techie buzzword. In many contexts it is abused as a synonym for the term *library*. Instead, we believe that a framework is something different, following the definition given in the famous design patterns book by the gang of four [31].

We use the term *library* when the root of the call graph is on the client side and she invokes the library function sparsely to obtain some concrete functionality. Instead, a *framework* stands in the root of the call graph and the client code is called through extension hooks in the framework, following the “Hollywood principle” — “Don’t call us, we’ll call you.”.

Because the Psychosynth system is layered, one can just use the bottom layers as a library, or rely on the upper layers that properly use inversion of control like a framework.

2.4.2 *The layered architecture*

At the current stage, GNU Psychosynth implements a layered architecture [32]. This intends to promote a more decoupled design, as calls between modules are only allowed from top down.

Also, the library has many features, many of which some users may not need. This layered approach could allow a potential user to avoid any special overhead when he is only using some bottom layers. Still, note that the library is now compiled with all layers bundled in the same shared-object

file, so this is not a fact now. Because the heavy redesign ongoing during this project, we shall postpone that until the late development stages when the layer interactions are clear and stable.

Figure 2.2 represents the current layered architecture. Lets make a bit more in-depth discussion of each layer.




Figure 2.2.: The Psychosynth layered architecture.

2.4.2.1 *The base layer*

The base layer includes basic facilities that may be reusable in any other part of the application. Some of the most relevant examples:

CONFIGURATION PERSISTENCE Classes for representing program settings in a hierarchical manner. This configuration system can use interchangeable backends and has an observable interface.²⁵

²⁵ We use quite often the term *observable interface* which is rare in the literature. By this, we mean that it provides signals, listeners or other event mechanisms, instances of the *observer* design pattern[31].

In fact, we do not recommend using this module in the core of the intermediate layers of the library because it can cause unexpected coupling, but this is not too clear and we keep it here for historical reasons.

IMPLEMENTATION OF DESIGN PATTERNS While the term *design pattern* means reusable design structure, not reusable code, language abstractions can make them implementable as code in some cases. Andrei Alexandrescu proves this point for C++ in [33]. Thus, we provide implementations, quite often similar to Alexandrescu's, to various recurring design patterns such as *factory* or *singleton*. Some implementations are not inspired in Alexandrescu's, like the generic *composite*.

COMMAND LINE ARGUMENT PARSING While we have considered moving to Boost's Program Options library, our own implementation have different trade-offs and is rather extensible.

LOGGING SYSTEM A hierarchical and multi backend logging system for registering messages from other modules. It should be used instead of direct output to `std::cout/cerr` in all the code.

FILE MANAGEMENT TOOLS That ease the task of finding resources in the hard-drive and can cache results.

Some other minor classes and tools are excluded from this list. During the development of the project we will drop in this layer classes that feel interesting at any abstraction level.

2.4.2.2 *The synth layer*

This layer contains classes for the *static* construction of synthesisers and sound processing. The audio input and output facilities are considered to be in this layer, and as well audio processing data structures —like ring buffers, multi channel buffersm, etc.—, basic implementations of filters, oscillators and audio scalers.

By static, we mean that this code does not provide any dynamic routing facilities, instead, the programmer is in charge to assign buffers and call the processing elements manually.

Requisites 1 to 3 should be implemented here. Non functional requisites 26 to 27 are specially relevant in this layer too.

2.4.2.3 *The node layer*

This layer provides the facilities for the *dynamic* construction of synthesisers. It includes the mechanisms for describing and executing the modular synthesis graph with the signal flow and so on. Figure 2.3 represents the main concepts behind the current design. Ports are considered as “signal ports” using the terminology in requirement 6 — “control ports” are similar to “parameters”, but parameters are not a precise model of “control ports” as they can not be routed and are intended for communication between the client user interface code and the audio thread state.

The communication system used to propagate values between the audio processing thread and the client thread is represented in figure 2.4. Values are copied to and from an intermediate channel between the audio processing blocks.




Figure 2.3.: Representation of the node graph as in Psychosynth 0.1.7. Input ports are represented as pentagons, output ports as rombus and parameteters as squares.




Figure 2.4.: Communication between the audio and user interface thread as in Psychosynth 0.1.7

Requisites 4 to 11 should be implemented in this layer. A heavy redesign of its API and many of its internal implementation is to be expected for that to be accomplished.

2.4.2.4 *The world layer and the Model View Controller architecture*

This layer simplifies the interface exposed to the previous layer and makes it *observable*. This is fundamental for the Model View Controller that the system implements. Figure

2.5 represents this architectural style. On the following, we can refer to this observable interface abstracting the synthesis engine as *the model*.

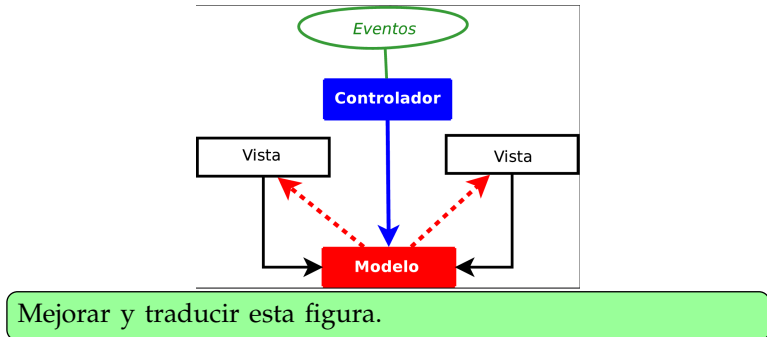


Figure 2.5.: The Model View Controller architectural style. Dashed arrows represent indirect invocation — i.e. via the *observer* design pattern— and normal lines represent normal method calls and data access.

Several views can coexist independently—for example, a GUI user interface and a web client—, that get updated whenever a value has changed in the model. They register themselves on the model at the beginning and then become passive entities that get called by the model. The model changes when controllers invokes methods on it, several controllers can coexist too. Usually, models and views come in pairs. For example, a GUI view has an associated controller that triggers the manipulation of the model in the eventuality of certain user actions like clicking a button; in this case the representation (the buttons) and the action (clicking it) are strongly related, but this is not necessarily true for other situations.

This layer also abstracts the graph interconnection mechanism using the *strategy* design pattern. Concretely, dynamic patching is implemented here and the interface exposed in this layer hides the manual node interconnection mechanisms but provides observability for topological changes.

This layer should implement requisite 16.

2.4.2.5 *The application framework layer and the view and controller layers*

There is a thin layer, instance of the *facade* pattern, called *app*, that was hidden for simplicity in figure 2.2 representing the layered architecture. It sits on top of the *world* layer and is in charge of initialising the *world* and defines a configuration tree, using the facilities in the base layer, and setups the machinery using the observability of the configuration system to keep coherence between the user preferences and the status of the synthesis model — for example, if the “*psynth.output*” setting is changed, it automatically creates the appropriate output object instance, sets its values and substitutes the output node in the synthesis graph. This layer also sets up the command line argument parsing and installs some common audio setting arguments in the parser. This layer is where Psychosynth becomes a framework at its most pure level, as it offers a *psynth_app* class whose *run* method should be the only call in the program *main*, and in turn delegates the heavy work to user code that is hooked as method overrides of that same class.

Orthogonal to this layer and sitting also on top of the *world* layer the networking layer offers a set of views and

controllers²⁶ that can be used to create the shared environment described in requirement 22 and thus allowing collaborative performances. This is an example of the value of the MVC architecture: because views and controllers are orthogonal, the user interface does not need to know about the presence of this to function properly; one could develop a new experimental and cool UI and it would automatically be able to work with third party clients over the network, even potentially using a different user interface.

On top of all this there is in the current version of Psychosynth the code of the 3D user interface and small and simple command line tools intended to be used as GUI-less network capable clients and servers. But all this code is not part of the framework, as it wants to be user interface agnostic, so we will not further describe that code.

2.5 PROJECT PLANNING AND METHODOLOGY

2.5.1 *Rationale — A critique on software engineering*

Lo se, soy un críticón enfermizo. He acabado convirtiendo esta sección en una crítica (indirecta y relativamente elegante, espero) a la concepción de la ingeniería del software que se da en la facultad y que se hace a menudo de los PFC. Si pensáis que esta crítica es demasiado larga o inapropiada puedo matizarla o eliminarla :D No obstante, en su defensa la considero apropiada ya que la elección de una metodología debe ser racional y como tanto debe incluir un análisis crítico de las alternativas.

²⁶ A primitive implementation of such at this stage of the development.

Choosing a well known software engineering process is considered one of the first steps to be taken in a final master thesis project. In our school we study with most detail the Cascade Process and the Unified Rational Process.

Those development processes propose a fordist software production model, targeted at huge development teams and the development of stable code bases in non innovative well defined fields. Martin Fowler makes a great point [34] criticising the often repeated argument stating that software engineering is like any other engineering where creative analysis and design is only the first step and thus coding is analogous to construction. As he says, the construction is done by the compiler and people involved in programming are actually doing an intellectual and creative work too — in computing, any systematic task can and must be automated. The *programming = construction* metaphor is alienating for the programmer, who is completely excluded from the task of criticising and improving the software design, and thus this metaphor often leads, in the end, to bad software.

Moreover, fordist development models take risk control and client requirements satisfaction as most important factors. Because we are in an academic environment, there are two more important factors: the pedagogical value of the project —this is, that the student involved takes the risk of exploring the unknown by himself— and the research value —this is, that the student involved takes the risk of exploring the unknown by humanity.

Of course, this is neither a pure research project, so we can not substitute a software development project by the scientific method. But we can choose a more dynamic methodology that includes *falsification* in one way or the other; *agile*

methodologies²⁷ propose many alternatives that could be valid for a master thesis project.

Still, these methodologies are, we believe, inadequate for this concrete project. The main reason is that this project is developed by only one person. Most methodologies, specially agile ones, put emphasis on the developer communication methods and collective decision making, so they are often inadequate and too constraining and time consuming for an unipersonal team. The Personal Software Process [35] proposes a methodology that is specially targeted at personal software developed by engineering students. Sadly, we are not very familiar with it—and do not have enough time to make that happen within the time constraints of the project—and it seems too be too specific and time consuming in its time tracking proposal.

Because we still believe that some rational planning and methodology is needed, we propose in the following a defined but unconstrained methodology that is specially tailored for our circumstances, capturing the most common elements in other software processes.

2.5.2 *An iterative development model*

Because of the size and complexity of the project, we should not consider developing it all at once. Moreover, the layered architecture of the starting code base and the variety of requirements that we want to satisfy favour an iterative development.

²⁷ <http://agilemanifesto.org/>

For all this, we want to split the development in iterations. Each iteration is composed by the following phases: *design*, *implementation*, *verification* and *integration*. Each iteration shall be assigned a set of requirements from the specification in section ?? that are to be satisfied after the successful accomplishment of that iteration.

2.5.2.1 *The design phase*

In the design phase we shall define the API that we would like the current subsystem have. Because we are developing a library and framework with a public interface, the design phase is specially relevant shall be done with care.

We do not enforce a particular method for documenting the design as different paradigms favour different documentation means. For example, in the first iteration we will develop a library heavily based on metaprogramming, where UML does not fit very naturally. Still, the documentation should include rationale explaining why the design decision lead to the satisfaction of the requirements that we want to satisfy after this development iteration. Also, it may be found that a requirement may be impossible to satisfy on the current iteration or that this requirement is to be better integrated in some other iteration. The developer is free to reassign that requisite for later iteration properly documenting this as a post-analysis plan fix.

What we do enforce is that all the API is documented with Doxygen for the sake of completeness of the reference manual that you should receive along with this document.

2.5.2.2 *The implementation phase*

During the implementation phase the code implementing the design should be written. It is possible and even recommended to modify the design during this phase as inconsistencies and fundamental problems are found. Sometimes, this may even start as soon as design, specially when it is unclear the properties that such API should have and some “exploratory programming” is needed. This fact may or many not be documented at the beginning of the section — even though an API may be designed through an inductive empirical process, a deductive rational description may be more useful for its clear understanding.

2.5.2.3 *The verification phase*

In the verification phase we perform *unit tests* on the most important parts of the system. No iteration should be considered finished unless proper unit tests are written and satisfied for its core components. For writing such tests the Boost Unit Testing Framework should be used.

When some elements are considered relevant to performance requirements, *performance tests* should be included. While we do not enforce a specific performance testing technique here, the tests should be reproducible and automatable whenever possible.

2.5.2.4 *The integration phase*

When a subsystem is added and it is to replace an existing subsystem in the project, the older code should be removed and the layers on top modified such that they use the new

code. This might even be considered part of the verification, as older tests working on the upper layers should be checked to be working still.

Informal integration tests should be done on the final user interface to make sure that the older properties are preserved. Note that in most cases, we do not recommend to lose time editing the old user interface such that the new features in the framework are exposed to the user. Of course, that the new features are usable is the final objective, but as it was justified in 1.2.1, a completely new user interface will be developed as part of a future project.

2.5.2.5 *Recursive decomposition of iterations*

In practice, some of the expected requirements to be satisfied may be found orthogonal or maybe too big to be addressed at one. It is thus allowed to recursively decompose an iteration in sub-iterations when a first evaluation during the design phase suggests that.

2.5.3 *A project plan*

2.5.3.1 *First iteration: A metaprogramming based sound processing foundational library*

This iteration is here to deeply re-design the core data structures using the latest techniques in C++. This requires special research and performance requirements deeply rely on the success of this iteration.

Requirements 1 to 3 and 26 to 27 should be satisfied for its success.

Estimated time cost: 6 weeks.

2.5.3.2 *Second iteration: Redesign of the node layer for hierarchy and polyphony*

The node layer requires a redesign if we want to satisfy all our purposes. Polyphony and hierarchy would be specially tricky to implement directly on top of the current code base. A special evaluation of how the new design interacts with the MVC architecture and networking is required. The world layer may be affected too. All the design changes should be implemented too.

Requirements 4 to 11 should be satisfied for its success. Requirement 23 may be considered for its implementation in this iteration too.

Estimated time cost: 6 weeks.

2.5.3.3 *Third iteration: Dynamic loading of nodes*

In this iteration the plugin system is to be developed.

Requirements 12 to 15 should be satisfied for its success.

Estimated time cost: 4 weeks.

2.5.3.4 *Fourth iteration: Adding MIDI and synchronisation*

Synchronisation and MIDI support is one of the most important features and it is also one of the features we know the least about, thus, we should put special care on research and design. This will affect the node and world layers mostly.

Requirements 17 to 22 should be satisfied for its success.

Estimated time cost: 8 weeks.

2.5.3.5 *Post mortem analysis*

After the conclusion of all the previous iterations, we should write a conclusive report and evaluation of its success. Also, we should prepare a final presentation for its evaluation.

3

A GENERIC SOUND PROCESSING LIBRARY

Numbers it is. All music
when you come to think.
Two multiplied by two
divided by half is twice
one. Vibrations: chords
those are. One plus two
plus six is seven. Do
anything you like with
figures juggling. Always
find out this equal to that.
Symmetry under a
cemetery wall. He doesn't
see my mourning.
Callous: all for his own
gut. Musemathematics.
And you think you're
listening to the etherial.
But suppose you said it
like: Martha, seven times
nine minus x is thirtyfive
thousand. Fall quite flat.
It's on account of the
sounds it is.

Ulysses
JAMES JOYCE

Este capítulo en general necesita más dibujos.

3.1 ANALYSIS

Requirements 1 to 3 and 26 to 27 refer to the second layer of our system — in a bottom-up approach. The most crucial question here: how do we represent a sound in a computer? Then, a new question arises: how do we get the sound to the speakers? The later question has a trivial answer — use whatever API your operating system exposes for delivering sound to the soundcard — but the first question is still to be answered. Actually, the solution to this first question mostly subsumes the issue of how to interface with these external interfaces, thus, shall we debate it with care.

3.1.1 *Factors of sound representation*

Me he dado cuenta de que esta sección introduce muchas definiciones. ¿Tal vez utilizar un estilo más formal usando el paquete theorem para las definiciones ayudaría a usar el documento como referencia?

A sound signal is a longitudinal wave of air in motion. We can analogically record the *proximal stimuli* — i.e. the physical stimuli leading to the subjective act of perception [36] — sound by measuring the successive oscillating values of air pressure in the observer's point in space. Note that an static air pressure value can not be perceived and the sensation of sound is caused by the relative oscillation of this measure. The *amplitude* of this change is associated to

our perception of *loudness*, the frequency of this oscillation mostly logarithmically determines our perception of pitch. We phrased this conditionally because these two variables are actually interrelated and our actual subjective perception of loudness might vary with pitch and otherwise [37].

Most of the time, we represent the relative air pressure value as a voltage, that varies at some range — p.e. $[-5, 5]V$. This is an analogous signal that we have to discretise somehow in order to manipulate it computationally.

3.1.1.1 *Temporal quantisation*

Temporal quantisation relates to how many times per second do we store the current voltage or air pressure value. The value of the signal between to equally spaced in time samples is unknown, but we can use some interpolation method to *upsample* a signal — i.e. to figure out what is between two samples. Most of the time we refer to the *sampling rate*, in hertz, as the frequency of the temporal quantisation.

We know from *Nyquist-Shannon sampling theorem* that perfect reconstruction of a signal is possible when the sampling frequency is greater than twice the maximum frequency of the signal being sampled, or equivalently, when the *Nyquist frequency* (half the sample rate) exceeds the highest frequency of the signal being sampled. Because the hearing range in most human beings is 20 Hz–20 kHz, audio compact discs use a 44.1 kHz sampling rate. Other popular rates in audio production are 48 kHz, 96 kHz and 192 kHz. Sampling rates bellow 44.1 kHz are used also in old computer games that were limited by the computing power and low bandwidth systems such as telephone, where low cost and proper un-

derstanding of human speech is more important than audio fidelity.

The sound representation mechanism itself does not vary with the sampling rate, and thus supporting various rates depends more on the implementation of the signal processing units and the overall performance of the system, with the CPU being able to process so many samples per second being the biggest constraint.

3.1.1.2 *Spatial quantisation*

Spatial quantisation determines how many possible values can a sample take in our finite and discrete scale. In a computer system, this is determined by the size in bits of the underlying type used to store the samples. An audio CD uses a *bitdepth* of 16 bit with samples that can take 65.536 possible values while professional audio uses 24 bit or 32 bit samples. We can even see systems using 64 bit samples during the processing to avoid accumulative rounding problems due to heavy arithmetic. The *dynamic range* of a signal with Q -bits quantisation is:

$$\text{DR}_{\text{ADC}} = 20 \times \log_{10}(2^Q) = (6.02 \cdot Q) \text{ dB} \quad (3.1)$$

The maximum *signal-to-noise* ratio for such a system is:

$$\text{SNR}_{\text{ADC}} = (1.76 + 6.02 \cdot Q) \text{ dB} \quad (3.2)$$

Most analog systems are said to have a dynamic range of around 80 dB [38]. Digital audio CD have a theoretical dynamic range of 96 dB — actual value is around 90 dB due to processing. Human hearing pain level is at 135 dB but actually prolonged exposure to such loud sound can

cause damage. A loud rock concert is around 120 dB and a classical music performance is at 110 dB [39], thus requiring bitdepth of at least 24 bit (theoretical dynamic range of 144 dB) for perfect fidelity.

There are some other aspects related to representation of samples in a computer, such as the *signedness* of the underlying type. Signed types are usually considered more convenient for audio signals as 0 can be easily recognised as the still no-sound value simplifying computations. Another important factor is whether we use *fixed point* or *floating point* arithmetic. While fixed point is used in low-cost DSP hardware, floating point is the most common representation in current audio software as nowadays processors are optimised for SIMD¹ floating point arithmetic. Moreover, the algorithms implementation is much harder to encode because products produce greater values and there are many ways on how to account the carry. Actually, even while the actual bitdepth (the bit for the mantissa) of a 32 bit floating point is the same of a 24 bit fixed point, then a 32 bit fixed point will have a quite lower *quantization error*, but the dynamic range and SNR of a floating point is much higher because the values are spaced logarithmically over a huge range [40]. Another factor is the *endianess* of fixed point values but is relevant only when interfacing with file formats and output devices.

3.1.1.3 *Channel space*

Because our hearing system is dicotomically symmetric, audio engineers discovered that much better fidelity can be

¹ Single Instruction Multiple Data, as supported by MMX, 3D Now! and SSE extensions in Intel and AMD chips

achieved by reproducing the sound with some differences from two separate loudspeakers. This is the well-known *stereophonic* sound, commonly named just *stereo*.

For representing such a signal, two different streams of information are needed for the left and right channels. Moreover, nowadays *quadraphonic*, and *surround* sound with varying numbers of channels up to 20.2 are used in different systems.

We call *channel space* to the set of semantic channels we use in some sort of audio representation — p.e. stereo audio has channel space with *left* and *right* elements. We use the term *frame* to call a set of samples coincident in time, this is, the samples of the various channels at a given time point. Thus, we will use most of the time the more accurate term *frame rate*.

This rises the problem on how to linearise the multi channel data. The most common mechanism in domestic hardware is by *interleaving* the samples of different channels, this is, by storing the frames sequentially. However, high-end hardware often accepts data in non-interleaved form where the samples of each channel is stored in a separate sequence. In this document, we borrow from the image processing world the term *planar* to refer to non-interleaved data. Software doing a lot of processing of the audio signal often chooses this representation as it is easier to scale to varying number of channels and split the signal to do per-channel filtering. Figure 3.1 compares visually the interleaved and planar formats.

Another issue is the order in which the data from different semantic channels is stored. We call a *channel layout* a bijection $L : C \rightarrow \mathbb{Z}_{\|C\|}$, where C is a given *channel space*.





Figure 3.1.: Multi-channel data in planar (a) and interleaved (b) form.

For example, the mapping $\{left \mapsto 0, right \mapsto 1\}$ is a common layout for stereo sound, but $\{left \mapsto 1, right \mapsto 0\}$ is sometimes used too.

3.1.2 Common solutions

As we have already noticed, 32 bit floating point sound with planar left-right layout the most common in software of our kind during internal processing. As most of this software is written in C, a simple `float**` does the job. This was, actually, the internal representation used in GNU Psychosynth in versions prior to 0.2.0, wrapped in the `audio_buffer` class.

However, this design starts to wobble whenever one has to interface with some other library or hardware using a different format. Thus, the `audio_buffer` class provided different `interleave_*` and `deinterleave_*`, where the asterisk can be substituted by different sample formats like `s16` or `s32` (fixed point signed 16 bit and 32 bit respectively). This is very inconvenient because, as we have seen through this

section, many orthogonal factor affect audio representation inducing a combinatorial explosion of format conversion functions. Take a look at the 64 different read and write functions in the `pcm.c` file of the `LibSndfile`² library.

This is a maintenance hell, but using the common means for abstracting orthogonal behaviour variability, i.e. dynamic polymorphism, is simply not an option in any audio software which supports real-time operation.

3.1.3 A generic approach: *Boost.GIL*

However, there is a piece of software that proved that this issue can be solved in C++ using static polymorphism. This is the Generic Image Library³ which was developed by Ludovic Cortés et. Al inside Adobe Software Technology Lab that was later include inside the Boost library distribution.

While sound and image manipulation are quite different, specially from the psycho-perceptive point of view, they are both a signal processing problem and thus share a lot in the representational issue. By realising of a proper conceptual mapping between both worlds (table 3.1), most of the library design and even quite a lot of code of *Boost.GIL* can be reused to build a unique state-of-the-art sound processing library that addresses the aforementioned issues in an orthogonal generic manner while maintaining near-optimal performance.

An *image* is bidimensional matrix of *pixels*, that capture the properties of light electromagnetic waveform at those discrete points. Each pixel, however, is decomposed in

² <http://www.mega-nerd.com/libsndfile/>

³ <http://stlab.adobe.com/gil>

Boost.GIL	Psynth.Sound
Channel	Sample
Color	Channel
Color Space	Channel Space
Color Layout	Channel Layout
Pixel	Frame
View	Range
Image	Buffer

Table 3.1.: Terminology map from `boost::gil` to `psynth::sound`

several *colors* that, for example, capture the intensity in the red, green and blue sensors of a CCD camera. As there are different ways of decomposing an audio frame (p.e, stereo, surround, etc.), there are different ways of decomposing a pixel into several values, known as the *color space* (p.e, RGB, CMYK, YUV, etc.). Boost.GIL uses the term *channel* to name the individual value of one those color components.

In our audio framework, a *buffer* is unidimensional array of *frames* that represent a sound or part of a sound — sound is continuous and thus we usually process it in chunks. The reader might note that the the data in a buffer being arranged along the *time* dimension while the dimensions of an image represent *physical space* makes these entities completely different from the processing point of view. However, they share most representation problems, with sound representation being actually a sub-problem of image representation, as we have one dimension less. The samples in a series of audio frames can be stored in an interleaved or planar fashion as happens with the channels of a pixel.

Also, both channels and samples can vary in signedness, fixed/floating point, bitdepth, etc.

Those already familiar with Boost.GIL can thus already understand easily our Psynth.Sound module design and implementation that we are to describe in the following section.

3.2 DESIGN

3.2.1 *Core techniques*

The Boost.GIL and thus the Psynth.Sound modules design makes heavy use of static polymorphism and generic programming via C++ templates to achieve generality without runtime overhead. We are going to introduce advanced techniques used in generic programming for the reader unfamiliar with this programming paradigm.

3.2.1.1 *Concepts*

.

Concepts [41] are to generic programming what *interfaces* — pure abstract classes in C++ — are to object oriented programming: they specify the requirements on some type. However there, are few substantial differences. (1) While interfaces can only specify the method signatures of its instances, a concept can specify most syntactic constraints on a type, like the existence of free functions, operators, nested types, etc. (2) While dispatching through interfaces requires, at least, a dereference, addition and function call [42], when using concepts the concrete function to be executed can be

determined and even inlined at compile-time. (3) One can not declare that a type satisfies an interface separately from the type definition, but one can say that a type models a concept at any point of the program. (4) Thus, no primitive type defines any virtual interface, but one can turn any primitive type into an instance of any concept via a `concept_map`. (5) Actually, the syntactic properties defined by a concept its models may differ, but they are matched via the `concept_map`. In fact, C++ concepts are more similar to Haskell *type classes*, with `instance` doing the job of `concept_map` [43].

Concepts are an extension to the template mechanism to add type checking for it. In fact, checking and dispatching on requirements can be achieved with techniques like SFINAE (Substitution Failure Is Not an Error) [44]. Property (5) of our concepts can be simulated with *traits* [45]. However, both compiler errors and the code using templates without concepts is usually much more unreadable.

The proposal of adding concepts to the C++ language was rejected last year by the standardisation committee and thus we can not use them in our code. However, Boost.GIL is very influenced by Alexander Stepanov's deductive approach to computer programming using generic programming and modeling with concepts, that he elegantly describes in his master-piece "Elements of Programming" [46]. Actually Stepanov worked several years in Adobe where he held a course "Foundations of Programming" based on his book. Thus, the *modeling* of the library extensively uses concepts. Its implementation uses a limited form of concept checking via the Boost.ConceptCheck⁴ [47] library, however, enabling

4 http://www.boost.org/doc/libs/release/libs/concept_check/concept_check.htm

this library in release mode can affect performance and its syntax is quite more cumbersome than the concepts in the C++ standard proposal. For consistency with the Boost.GIL documentation we will use the concept syntax proposed in the proposal N2081 to the standardisation committee [48].

The following example defines a concept that is satisfied by every type that has an operator<:

```
concept LessThanComparable<typename T> {
    bool operator< (T, T);
};
```

This allows us to write a generic function that depends on the existence of a less-than comparator for the parametrised type:

```
template<LessThanComparable T>
const T& min (const T& x, const T& y) {
    return x < y? x : y;
}
```

An alternative syntax for specifying that T must satisfy the LessThanComparable concept is the where clause:

```
template<typename T>
    where LessThanComparable<T>
const T& min (const T& x, const T& y) ...
```

In fact, this is the only valid syntax when the concept affects multiple types. Also, the where clause can be used inside concept definitions to provide specialisation.

Specifying that a type models a concept is done with the `concept_map` device. If the type naturally models the concept, we can just use:

```
concept_map LessThanComparable<int> {}
```

Note that these trivial concept mappings can be avoided by using the `auto` keyword in front of the concept keyword in the concept definition. However, it might happen that a type requires some wrapping to satisfy the concept. We can do this in the concept map definition itself.

```
concept_map LessThanComparable<char*> {
    bool operator< (char* a, char* b) {
        return strcmp (a, b) < 0;
    }
}
```

Note that this last piece of code is an example of a bad usage of concept maps, as this specialises the mapping for pointers changing the expected semantics.

This should suffice as an introduction to concepts in order to understand the concept definitions that we will later show when modelling our system. A more detailed view can be read in the cited bibliography, with [41] being the most updated and useful from a programmer point of view.

3.2.1.2 *Metaprogramming*

The C++ template system is Turing complete [49], thus it can be used to perform any computation at *compile time*. This was first noted in 1994 by Erwin Unruh who, in the middle of a C++ standardisation committee, wrote a template metaprogram that outputted the first N prime numbers on the console using compiler errors [50]. Even though this might seem just a crazy puzzle game, it can be used in practise and actually new Boost libraries use it extensively. A very gentle introduction to template metaprogramming can be found in [33], where Alexandrescu uses them to instantiate design patterns as generic C++ libraries. A deeper reference is Abra-

ham's [51], which focuses on the Boost Metaprogramming Library⁵ and introduces the usage of metaprogramming for building Embedded Domain Specific Languages (EDSL) in C++. This Boost.MPL, providing reusable meta data structures and algorithms, is the de-facto standard library for template metaprogramming⁶ and we will use it in our implementation.

Template metaprogramming is possible thanks to *partial template specialisation*, that allows giving an alternate definition for a pattern matched subset of its possible parameter values. A *metafunction* is thus just a template class or struct with a public member that holds the result of the function. It is up to the programmer to choose the naming convention for the result members of the metafunctions, in the following, we will use Abraham's style calling type for result values that are a type, and value for integral values. Listing 3.1 illustrates how can we write and use a metafunction for computing the n -th Fibonacci number.

The program returns the forty-second Fibonacci value. However, it will take no time to execute, because the number is computed at compile time. We use recursion to define the metafunction for the general case and the specialise for the base cases.

If we consider the template system as a meta-language on its own, we should describe its most outstanding semantic properties. It is a pure functional programming language, because variables are immutable. It is lexically scoped. It supports both lazy and strict evaluation, depending on whether we choose to access the nested type result name at call site or value usage type. When we look at the

⁵ The Boost.MPL: www.boost.org/doc/libs/release/libs/mpl

⁶ It is often called "the STL of template metaprogramming".

Listing 3.1: Metaprogram for computing the Nth Fibonacci number

```

template <int N>
struct fib {
    enum {
        value = fib<N-1>::value + fib<N-2>::value;
    };
};

template <>
struct fib <0> {
    enum { value = 0 };
};

template <>
struct fib <1> {
    enum { value = 1 };
};

int main () {
    return fib<42>::value;
}

```

meta type system, we find three meta types: types (which are duck-typed records), integrals (e.g. `int`, `char`, `bool` ...) and meta-functions (i.e. templates).

The fact that records are duck typed but integrals and metafunctions cause several inconveniences in practice, specially when dealing with the later. For example, in the absence of template aliases, returning a metafunction produced by another function requires defining a nested struct that inherits from the actual synthesised value. Also, the

Listing 3.2: Integral constant nullary metafunction wrapper.

```
template <typename T, T V>
struct integral_c
{
    BOOST_STATIC_CONSTANT(T, value = V);
    typedef integral_c <T, V> type;
};
```

template signature should be specified on a template parameter expecting a template.

In order to simplify our meta type system we shall wrap constants in a type like on listing 3.2.

There are a couple of issues regarding this definition worth explaining. First, the `BOOST_STATIC_CONSTANT` macro is used to define a constant. Internally, it will try to use `enum` or any other mechanism available to actually define the constant such that the compiler is not tempted to allocate static memory for the constant. Second, the `typedef` referring to itself turns a constant value into a self returning nullary meta-function. This can be very convenient because, for example, it allows using `value::type::value` always on the value usage point, allowing the caller or producer of the value to choose whether he wants to evaluate the value lazily.

Because we just wrapped values into a type, we can simplify our conventional definition of *metafunction*: a metafunction is any type — template or not — that has a nested type called `type`.

Now we should also turn metafunctions into first class entities of the meta-language. We just add a new level of

Listing 3.3: Metafunction class for computing Fibonacci numbers. We suppose that the previous `fib` definition uses `integral_c` to wrap its parameters and return types.

```

struct fib_class {
    template <class N>
    struct apply : public fib<N> {};
};

int maint ()
{
    return fib_class :: apply<integral_c<int, 42>>
        :: type :: value;
}

```

indirection and define a *metafunction class* as a type with a nested template metafunction called `apply`. The example in listing 3.3 also illustrates the metafunction forwarding technique when defining the nested `apply` metafunction by inheriting from `fib`.

Using this convention the MPL library defines many useful high order metafunctions that take metafunction classes as input, like `mpl::fold` and `mpl::transform`. Note that it is not needed to define metafunction classes for all our metafunctions, instead, we shall convert them when needed using the `mpl::quoteN` functions and the `mpl::lambda` facility.

3.2.2 Core concepts

We are now ready to understand the main design and implementation techniques used in our generic library. Because the library is *generic*, in the sense of generic programming, most algorithms and data structures are parametrised such that they can be instantiated with any concrete type modelling some concepts as we suggested in section 3.2.1.1. Thus, traditional modelling techniques like the Unified Modelling Language are not useful since they are intended for object oriented design.

We are going to follow the following methodology for describing the library. First, we will name a concept and give a brief description of its purpose. Then, we will define the concept using the notation described in section 3.2.1.1 and finally we will enumerate and describe some models for such concept.

For brevity, we will omit basic concepts such as `CopyConstructible`, `Regular`, `Metafunction`, etc. Their complete definition should be evident and an interested reader can find most of them in [46].

3.2.2.1 *ChannelSpaceConcept*

A channel space is a sequence of whose elements channel tags (empty types giving a name)

```
concept ChannelSpaceConcept<
    MPLRandomAccessSequence Cs>
{
};
```

Some example models include `stereo_space` or `surround_space`. An example on how a user of the library can define his own channel space follows.

```
struct left_channel {};
struct right_channel {};

typedef mpl::vector<left_channel, right_channel>
    stereo_space;
```

A related trivial concept is `ChannelSpaceCompatibleConcept`. Two channel spaces are compatible if they are the same. In fact, this leaks the underlying MPL sequence type used in the channel space through the abstraction, because spaces with the same set of semantic channels might be found incompatible, but it suffices in practise.

3.2.2.2 *SampleConcept*

A *sample* is the type we use to represent the amplitude of a channel at certain point in time.

```
concept SampleConcept<typename T> :
    EqualityComparable<T> {
    typename value_type      = T;
    // use sample_traits<T>::value_type to access it
    typename reference      = T&;
    // use sample_traits<T>::reference to access it
    typename pointer        = T*;
    // use sample_traits<T>::pointer to access it
    typename const_reference = const T&;
    // use sample_traits<T>::const_reference to access it
    typename const_pointer  = const T*;
    // use sample_traits<T>::const_pointer to access it
    static const bool is_mutable;
    // use sample_traits<T>::is_mutable to access it
```

```

static T min_value();
// use sample_traits<T>::min_value to access it
static T zero_value();
// use sample_traits<T>::zero_value to access it
static T max_value();
// use sample_traits<T>::min_value to access it
};

```

Built-in scalar types like `char`, `int` or `float` model `SampleConcept` by default.

The `scoped_sample_value<Type, Min, Max, Zero>` template class models the concept whenever `Type` is a scalar type and `Min`, `Zero` and `Max` satisfy:

$$Min < Zero < Max \wedge \forall x \in Type, x + Zero = x \quad (3.3)$$

Note that, in order to avoid the limitation of floating point values not being usable as template arguments, `Min`, `Zero` and `Max` should be a type with a static method `apply()` that returns the actual value. It should be used to constraint the “clipping thresholds” of floating point types. For example, the `bits32sf` model defined as:

```

scoped_sample<float,
              float_minus_one ,
              float_zero ,
              float_one >;

```

User defined types should specialise `sample_traits` to map the concept.

Related trivial concepts are `MutableSampleConcept` and `SampleValueConcept` (a sample that is also `Regular`).

3.2.2.3 *SampleConvertibleConcept*

Because casting does not suffice in most cases, one should override a `T sample_convert (U)` function for `U` to be convertible into `T`.

```
concept SampleConvertibleConcept<SampleConcept SrcSample,
    SampleValueConcept DstSample> {
    DstSample sample_convert (const SrcSample&);
};
```

The library provides overrides for `sample_convert` making most supplied sample types being convertible too.

3.2.2.4 *ChannelBaseConcept*

A *channel base* is a container of channel elements (such as samples, sample references or sample pointers).

The most common use of channel base is in the implementation of a frame, in which case the channel elements are sample values. The channel base concept, however, can be used in other scenarios. For example, a planar frame has samples that are not contiguous in memory. Its reference is a proxy class that uses a channel base whose elements are sample references. Its iterator uses a channel base whose elements are sample iterators.

```
concept ChannelBaseConcept<typename T> :
    CopyConstructible<T>, EqualityComparable<T>
{
    // a Psynth layout (the channel space and element permutation)
    typename layout;

    // The type of K-th element
    template <int K> struct kth_element_type;
```



```

where Metafunction<kth_element_type>;

// The result of at_c
template <int K>
struct kth_element_const_reference_type ;
where Metafunction<
    kth_element_const_reference_type >;

template <int K>
    kth_element_const_reference_type <T,K>::type at_c(T);

// Copy-constructible and equality comparable
// with other compatible channel bases
template <ChannelBaseConcept T2>
    where { ChannelBasesCompatibleConcept<T,T2> }
    T::T(T2);
template <ChannelBaseConcept T2>
    where { ChannelBasesCompatibleConcept<T,T2> }
    bool operator==(const T&, const T2&);
template <ChannelBaseConcept T2>
    where { ChannelBasesCompatibleConcept<T,T2> }
    bool operator!=(const T&, const T2&);
};

```

A channel base must have an associated layout (which consists of a channel space, as well as an ordering of the samples). There are two ways to index the elements of a channel base: A physical index corresponds to the way they are ordered in memory, and a semantic index corresponds to the way the elements are ordered in their channel space. For example, in the stereo channel space the elements are ordered as {left_channel, right_channel}. For a channel base with a RL-stereo layout, the first element in physical ordering is the right element, whereas the first semantic element is the red one. Models of ChannelBaseConcept are required to provide the at_c<K>(ChannelBase) function, which allows

for accessing the elements based on their physical order. `Psynth` provides a `semantic_at_c<K>(ChannelBase)` function (described later) which can operate on any model of `ChannelBaseConcept` and returns the corresponding semantic element.

The library provides an `homogeneous_channel_base` class that models the concept.

3.2.2.5 ... *Concept*

3.2.3 *Going dynamic*

3.2.4 *Input and Output module*

3.2.5 *Synthesis module*

3.3 VALIDATION

3.3.1 *Unit testing*

3.3.2 *Performance*

3.3.3 *Integration*

The module was first developed separately from the main development branch in a branch called `gil-import`. Once the previous tests were passed, the former signal representation classes and I/O code was removed from the code base. The upper layers — mainly the graph layer — was adapted to use the new library. Note that, given that we will

mostly rewrite the graph layer in the next iteration we tried to make minimal changes to get the project compile and run properly again.

The system was then *peer reviewed* by project collaborators, mainly by the maintainer of the Ubuntu/Trinux packages Aleksander Morgado⁷. After minor bugfixing, we agreed to make a new Psychosynth 0.2.0 release that included the new code described in this chapter and some other fixes and modifications developed alongside.

The official changelog briefing for this release includes:

- New audio processing and I/O subsystem based on template programming for generic yet efficient sound signals.
- The extreme latency when using ALSA bug seems to be fixed in some cases.
- No longer depend on libvorbis, libsndfile can now handle ogg and flac files too.
- No longer depend on libsigc++, using boost::signals which, which is a bit slower but neglibe and this simplifies the dependencies.
- The mouse wheel now scrolls in the object selector.
- The object selector no longer lets mouse clicks pass through.
- Backwards reproducing a sample works a bit better now too.
- Some new niceties in the framework base layer, including some experiments on applying the C3 class linearisation algorithm in raw C++.

⁷ <http://sigquit.wordpress.com>

- C++0x features are being used in the code. For GCC, this means version 4.5 shall be used. We doubt it will compile with any other compiler (maybe latest VS), but users are welcomed to try and report.
- For this same reason, Boost.Threads is no longer a dependency, we use STL threads instead.

3.4 CONCLUSIONS AND FUTURE WORK

BIBLIOGRAPHY

- [1] Sergi Jordà, Günter Geiger, Marcos Alonso, and Martin Kaltenbrunner. The Reactable: Exploring the synergy between live music performance and tabletop tangible interfaces. In *In Proceedings of the first international conference on "Tangible and Embedded Interaction"*, Baton, pages 139–146. ACM Press, 2007.
- [2] Paul Tingen. Autechre, recording electronica. *Sound on Sound*, 2004.
- [3] Tue Haste Andersen. Mixxx: towards novel dj interfaces. In *Proceedings of the 2003 conference on New interfaces for musical expression*, NIME '03, pages 30–35, Singapore, Singapore, 2003. National University of Singapore.
- [4] Robert A. Moog. Voltage-controlled electronic music modules. In *Audio Engineering Society Convention 16*, 10 1964.
- [5] M. Kaltenbrunner, G. Geiger, and S. Jordà. Dynamic patches for live musical performance. 2004.
- [6] Thor Magnusson and Enrike Hurtado Mendieta. The acoustic, the digital and the body: a survey on musical instruments. In *Proceedings of the 7th international conference on New interfaces for musical expression*, NIME '07, pages 94–99, New York, NY, USA, 2007. ACM.

- [7] MIDI Manufacturers Association Incorporated. *Complete MIDI 1.0 Detailed Specification*, 1999/2008.
- [8] Richard M. Stallman and Joshua Gay. *Free Software, Free Society: Selected Essays of Richard M. Stallman*. Free Software Foundation, June 2002.
- [9] Miller Puckette. Max at 17. In *Computer Music Journal*, volume 26 (4), pages 31–43, 10 2002.
- [10] Miller Puckette. Pure data: another integrated computer music environment. In *in Proceedings, International Computer Music Conference*, pages 37–41, 1996.
- [11] Stephane Letz, Nedko Arnaudov, and Romain Moret. What is new in jack 2? In *In Proceedings of the Linux Audio Conference 2009*, 2009.
- [12] Torben Hohn, Alexander Carot, and Christian Werner. Netjack: Remote music collaboration with electronic sequencers on the internet. In *In Proceedings of the Linux Audio Conference 2009*, 2009.
- [13] Tina Blaine and Sidney Fels. Contexts of collaborative musical experiences. In *NIME '03: Proceedings of the 2003 conference on New interfaces for musical expression*, pages 129–134, Singapore, Singapore, 2003. National University of Singapore.
- [14] James Patten, Ben Recht, and Hiroshi Ishii. Audiopad: a tag-based interface for musical performance. In *Proceedings of the 2002 conference on New interfaces for musical expression*, NIME '02, pages 1–6, Singapore, Singapore, 2002. National University of Singapore.
- [15] Tina Blaine. The jam-o-drum interactive music system: a study in interaction design. In *In DIS '00: Proceedings of the conference on Designing interactive systems*, pages

- 165–173. ACM Press, 2000.
- [16] R. Bencina, M. Kaltenbrunner, and S. Jordà. Improved topological fiducial tracking in the reactivision system. 2005.
 - [17] Matthew Wright Center and Matthew Wright. Open-sound control: State of the art 2003. In *In Proceedings of the 2003 Conference on New Interfaces for Musical Expression (NIME-03)*, pages 153–159, 2003.
 - [18] Martin Kaltenbrunner, Till Bovermann, Ross Bencina, and Enrico Costanza. Tuio: A protocol for table-top tangible user interfaces. In *In Proceedings of the 2 nd Interactive Sonification Workshop*, 2005.
 - [19] Yu Nishibori and Toshio Iwai. Tenori-on. In Norbert Schnell, Frédéric Bevilacqua, Michael J. Lyons, and Atau Tanaka, editors, *NIME*, pages 172–175. IRCAM - Centre Pompidou in collaboration with Sorbonne University, 2006.
 - [20] Gary P. Scavone and Perry R. Cook. Rtmidi, rtaudio, and a synthesis toolkit (stk) update. In *In Proceedings of the International Computer Music Conference*, 2005.
 - [21] R. Boulanger, editor. *The Csound Book: Perspectives in Software Synthesis, Sound Design, Signal Processing, and Programming*. The MIT Press, March 2000.
 - [22] James McCartney. Rethinking the computer music language: Supercollider. *Computer Music Journal*, 26(4):61–68, 2002.
 - [23] Ge Wang and Perry R. Cook. Chuck: a concurrent, on-the-fly audio programming language. In *Proc. ICMC*, pages 219–226, 2003.

- [24] Yann Orlarey, Dominique Fober, and Stephane Letz. *FAUST: an Efficient Functional Approach to DSP Programming*, pages 65–96. 2009.
- [25] Heinrich Taube. An Introduction to Common Music. *Computer Music Journal*, 21(1), 1997.
- [26] Xavier Amatriain, Pau Arumi, and David Garcia. Clam: a framework for efficient and rapid development of cross-platform audio applications. In *Proceedings of the 14th annual ACM international conference on Multimedia*, MULTIMEDIA '06, pages 951–954, New York, NY, USA, 2006. ACM.
- [27] Richard Stallman. Why you shouldn't use the Lesser GPL for your next library. *The GNU Project Philosophy webpage*: <http://www.gnu.org/licenses/why-not-lgpl.html>, last accessed 4 of January 2011., 1997.
- [28] Herb Sutter. Trip report: March 2010 iso c++ standards meeting. *Sutter's Mill Blog*: <http://herbsutter.com/2010/03/13/trip-report-march-2010-iso-c-standards-meeting/>, last accessed 04 of January 2011., 2010.
- [29] R. M. Stallman. *GNU Coding Standards*. FSF, <http://www.gnu.org/prep/standards/standards.html>, last visited January 5th 2011, 1992–2010.
- [30] Juan Pedro Bolívar Puente. GNU Psychosynth: Un sintetizador de audio por Software Libre. *Novática*, 195:47–50, September 2008.
- [31] Erich Gamma, Richard Helm, Ralph E. Johnson, and John Vlissides. *Design Patterns. Elements of Reusable Object-Oriented Software*. Addison-Wesley, March 1995.

- [32] David Garlan and Mary Shaw. An introduction to software architecture. Technical Report CMU-CS-94-166, Carnegie Mellon University, January 1994.
- [33] Andrei Alexandrescu. *Modern C++ design: generic programming and design patterns applied*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2001.
- [34] Martin Fowler. *Is design dead?*, pages 3–17. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2001.
- [35] Watts S. Humphrey. Using a defined and measured personal software process. *IEEE Software*, 13:77–88, 1996.
- [36] E. Goldstein. *Sensation and Perception*. Wadsworth Publishing Company, Belmont, California, 6th edition edition, 2001.
- [37] Harvery Fletcher. Loudness and pitch. *Bell Laboratories Record*, 13(5):130–137, January 1935.
- [38] Bruce Fries and Marty Fries. *Digital Audio Essentials*. O'Reilly Media, Inc., 2005.
- [39] Ars Ludwig. Music and the human hearing. Self published: <http://www.silcom.com/~aludwig/EARS.htm>, last accessed 03 of March 2011., 2009.
- [40] Steven W. Smith. *The scientist and engineer's guide to digital signal processing*. California Technical Publishing, San Diego, CA, USA, 2002.
- [41] Jaakko Järvi, Mat Marcus, and Jacob N. Smith. Programming with c++ concepts. *Science of Computer Programming*, 75(7):596 – 614, 2010. Generative Programming and Component Engineering (GPCE 2007).

- [42] Karel Driesen and Urs Hölzle. The direct cost of virtual function calls in c++. In *Proceedings of the 11th ACM SIGPLAN conference on Object-oriented programming, systems, languages, and applications, OOPSLA '96*, pages 306–323, New York, NY, USA, 1996. ACM.
- [43] Jean-Philippe Bernardy, Patrik Jansson, Marcin Zalewski, Sibylle Schupp, and Andreas Priesnitz. A comparison of c++ concepts and haskell type classes. In *Proceedings of the ACM SIGPLAN workshop on Generic programming, WGP '08*, pages 37–48, New York, NY, USA, 2008. ACM.
- [44] David Vandevoorde and Nicolai M. Josuttis. *C++ Templates: The Complete Guide*. Addison-Wesley Professional, November 2002.
- [45] Nathan C. Myers. Traits: a new and useful template technique. *C++ Report*, June 1995.
- [46] Alexander Stepanov and Paul McJones. *Elements of Programming*. Addison-Wesley Professional, 1 edition, June 2009.
- [47] Jeremy G. Siek and Andrew Lumsdaine. Concept checking: Binding parametric polymorphism in C++. In *Proceedings of the First Workshop on C++ Template Programming*, Erfurt, Germany, 2000.
- [48] Douglas Gregor and Bjarne Stroustrup. Concepts (revision 1). (N2081=06-0151), 10/2006 2006.
- [49] Todd L. Veldhuizen. C++ templates are turing complete. Technical report, Indiana University of Computer Science, 2003.
- [50] Erwin Unruh. Compile time prime number computation. ANSI X3J16-94-0075/ISO WG21-462. Available

online: <http://www.erwin-unruh.de/primorig.html>,
1994.

- [51] David Abrahams and Aleksey Gurtovoy. *C++ Template Metaprogramming: Concepts, Tools, and Techniques from Boost and Beyond (C++ in Depth Series)*. Addison-Wesley Professional, 2004.

Appendices



GLOSSARY

¿Añado un glosario recopilando las definiciones que doy de diferentes términos específicos del dominio del audio y para expandir y aclarar siglas y acrónimos? ¿O mejor paso? — JP



GNU GENERAL PUBLIC LICENSE

Version 3, 29 June 2007

Copyright © 2007 Free Software Foundation, Inc.
<http://fsf.org/>

Everyone is permitted to copy and distribute verbatim
copies of this
license document, but changing it is not allowed.

B.1 PREAMBLE

The GNU General Public License is a free, copyleft license for software and other kinds of works.

The licenses for most software and other practical works are designed to take away your freedom to share and change the works. By contrast, the GNU General Public License is intended to guarantee your freedom to share and change all versions of a program—to make sure it remains free software for all its users. We, the Free Software Foundation, use the GNU General Public License for most of our software;

it applies also to any other work released this way by its authors. You can apply it to your programs, too.

When we speak of free software, we are referring to freedom, not price. Our General Public Licenses are designed to make sure that you have the freedom to distribute copies of free software (and charge for them if you wish), that you receive source code or can get it if you want it, that you can change the software or use pieces of it in new free programs, and that you know you can do these things.

To protect your rights, we need to prevent others from denying you these rights or asking you to surrender the rights. Therefore, you have certain responsibilities if you distribute copies of the software, or if you modify it: responsibilities to respect the freedom of others.

For example, if you distribute copies of such a program, whether gratis or for a fee, you must pass on to the recipients the same freedoms that you received. You must make sure that they, too, receive or can get the source code. And you must show them these terms so they know their rights.

Developers that use the GNU GPL protect your rights with two steps: (1) assert copyright on the software, and (2) offer you this License giving you legal permission to copy, distribute and/or modify it.

For the developers' and authors' protection, the GPL clearly explains that there is no warranty for this free software. For both users' and authors' sake, the GPL requires that modified versions be marked as changed, so that their problems will not be attributed erroneously to authors of previous versions.

Some devices are designed to deny users access to install or run modified versions of the software inside them, al-

though the manufacturer can do so. This is fundamentally incompatible with the aim of protecting users' freedom to change the software. The systematic pattern of such abuse occurs in the area of products for individuals to use, which is precisely where it is most unacceptable. Therefore, we have designed this version of the GPL to prohibit the practice for those products. If such problems arise substantially in other domains, we stand ready to extend this provision to those domains in future versions of the GPL, as needed to protect the freedom of users.

Finally, every program is threatened constantly by software patents. States should not allow patents to restrict development and use of software on general-purpose computers, but in those that do, we wish to avoid the special danger that patents applied to a free program could make it effectively proprietary. To prevent this, the GPL assures that patents cannot be used to render the program non-free.

The precise terms and conditions for copying, distribution and modification follow.

B.2 TERMS AND CONDITIONS

o. Definitions.

"This License" refers to version 3 of the GNU General Public License.

"Copyright" also means copyright-like laws that apply to other kinds of works, such as semiconductor masks.

"The Program" refers to any copyrightable work licensed under this License. Each licensee is addressed

as “you”. “Licensees” and “recipients” may be individuals or organizations.

To “modify” a work means to copy from or adapt all or part of the work in a fashion requiring copyright permission, other than the making of an exact copy. The resulting work is called a “modified version” of the earlier work or a work “based on” the earlier work.

A “covered work” means either the unmodified Program or a work based on the Program.

To “propagate” a work means to do anything with it that, without permission, would make you directly or secondarily liable for infringement under applicable copyright law, except executing it on a computer or modifying a private copy. Propagation includes copying, distribution (with or without modification), making available to the public, and in some countries other activities as well.

To “convey” a work means any kind of propagation that enables other parties to make or receive copies. Mere interaction with a user through a computer network, with no transfer of a copy, is not conveying.

An interactive user interface displays “Appropriate Legal Notices” to the extent that it includes a convenient and prominently visible feature that (1) displays an appropriate copyright notice, and (2) tells the user that there is no warranty for the work (except to the extent that warranties are provided), that licensees may convey the work under this License, and how to view a copy of this License. If the interface presents a list of user commands or options, such as a menu, a prominent item in the list meets this criterion.

1. Source Code.

The “source code” for a work means the preferred form of the work for making modifications to it. “Object code” means any non-source form of a work.

A “Standard Interface” means an interface that either is an official standard defined by a recognized standards body, or, in the case of interfaces specified for a particular programming language, one that is widely used among developers working in that language.

The “System Libraries” of an executable work include anything, other than the work as a whole, that (a) is included in the normal form of packaging a Major Component, but which is not part of that Major Component, and (b) serves only to enable use of the work with that Major Component, or to implement a Standard Interface for which an implementation is available to the public in source code form. A “Major Component”, in this context, means a major essential component (kernel, window system, and so on) of the specific operating system (if any) on which the executable work runs, or a compiler used to produce the work, or an object code interpreter used to run it.

The “Corresponding Source” for a work in object code form means all the source code needed to generate, install, and (for an executable work) run the object code and to modify the work, including scripts to control those activities. However, it does not include the work’s System Libraries, or general-purpose tools or generally available free programs which are used unmodified in performing those activities but which are not part of the work. For example, Corresponding

Source includes interface definition files associated with source files for the work, and the source code for shared libraries and dynamically linked subprograms that the work is specifically designed to require, such as by intimate data communication or control flow between those subprograms and other parts of the work.

The Corresponding Source need not include anything that users can regenerate automatically from other parts of the Corresponding Source.

The Corresponding Source for a work in source code form is that same work.

2. Basic Permissions.

All rights granted under this License are granted for the term of copyright on the Program, and are irrevocable provided the stated conditions are met. This License explicitly affirms your unlimited permission to run the unmodified Program. The output from running a covered work is covered by this License only if the output, given its content, constitutes a covered work. This License acknowledges your rights of fair use or other equivalent, as provided by copyright law.

You may make, run and propagate covered works that you do not convey, without conditions so long as your license otherwise remains in force. You may convey covered works to others for the sole purpose of having them make modifications exclusively for you, or provide you with facilities for running those works, provided that you comply with the terms of this License in conveying all material for which you do not control copyright. Those thus making or running the

covered works for you must do so exclusively on your behalf, under your direction and control, on terms that prohibit them from making any copies of your copyrighted material outside their relationship with you.

Conveying under any other circumstances is permitted solely under the conditions stated below. Sublicensing is not allowed; section 10 makes it unnecessary.

3. Protecting Users' Legal Rights From Anti-Circumvention Law.

No covered work shall be deemed part of an effective technological measure under any applicable law fulfilling obligations under article 11 of the WIPO copyright treaty adopted on 20 December 1996, or similar laws prohibiting or restricting circumvention of such measures.

When you convey a covered work, you waive any legal power to forbid circumvention of technological measures to the extent such circumvention is effected by exercising rights under this License with respect to the covered work, and you disclaim any intention to limit operation or modification of the work as a means of enforcing, against the work's users, your or third parties' legal rights to forbid circumvention of technological measures.

4. Conveying Verbatim Copies.

You may convey verbatim copies of the Program's source code as you receive it, in any medium, provided that you conspicuously and appropriately publish on each copy an appropriate copyright notice; keep intact all notices stating that this License and any non-

permissive terms added in accord with section 7 apply to the code; keep intact all notices of the absence of any warranty; and give all recipients a copy of this License along with the Program.

You may charge any price or no price for each copy that you convey, and you may offer support or warranty protection for a fee.

5. Conveying Modified Source Versions.

You may convey a work based on the Program, or the modifications to produce it from the Program, in the form of source code under the terms of section 4, provided that you also meet all of these conditions:

- a) The work must carry prominent notices stating that you modified it, and giving a relevant date.
- b) The work must carry prominent notices stating that it is released under this License and any conditions added under section 7. This requirement modifies the requirement in section 4 to “keep intact all notices”.
- c) You must license the entire work, as a whole, under this License to anyone who comes into possession of a copy. This License will therefore apply, along with any applicable section 7 additional terms, to the whole of the work, and all its parts, regardless of how they are packaged. This License gives no permission to license the work in any other way, but it does not invalidate such permission if you have separately received it.
- d) If the work has interactive user interfaces, each must display Appropriate Legal Notices; however,

if the Program has interactive interfaces that do not display Appropriate Legal Notices, your work need not make them do so.

A compilation of a covered work with other separate and independent works, which are not by their nature extensions of the covered work, and which are not combined with it such as to form a larger program, in or on a volume of a storage or distribution medium, is called an “aggregate” if the compilation and its resulting copyright are not used to limit the access or legal rights of the compilation’s users beyond what the individual works permit. Inclusion of a covered work in an aggregate does not cause this License to apply to the other parts of the aggregate.

6. Conveying Non-Source Forms.

You may convey a covered work in object code form under the terms of sections 4 and 5, provided that you also convey the machine-readable Corresponding Source under the terms of this License, in one of these ways:

- a) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by the Corresponding Source fixed on a durable physical medium customarily used for software interchange.
- b) Convey the object code in, or embodied in, a physical product (including a physical distribution medium), accompanied by a written offer, valid for at least three years and valid for as long as you offer spare parts or customer support for that product model, to give anyone who possesses the

object code either (1) a copy of the Corresponding Source for all the software in the product that is covered by this License, on a durable physical medium customarily used for software interchange, for a price no more than your reasonable cost of physically performing this conveying of source, or (2) access to copy the Corresponding Source from a network server at no charge.

- c) Convey individual copies of the object code with a copy of the written offer to provide the Corresponding Source. This alternative is allowed only occasionally and noncommercially, and only if you received the object code with such an offer, in accord with subsection 6b.
- d) Convey the object code by offering access from a designated place (gratis or for a charge), and offer equivalent access to the Corresponding Source in the same way through the same place at no further charge. You need not require recipients to copy the Corresponding Source along with the object code. If the place to copy the object code is a network server, the Corresponding Source may be on a different server (operated by you or a third party) that supports equivalent copying facilities, provided you maintain clear directions next to the object code saying where to find the Corresponding Source. Regardless of what server hosts the Corresponding Source, you remain obligated to ensure that it is available for as long as needed to satisfy these requirements.
- e) Convey the object code using peer-to-peer transmission, provided you inform other peers where

the object code and Corresponding Source of the work are being offered to the general public at no charge under subsection 6d.

A separable portion of the object code, whose source code is excluded from the Corresponding Source as a System Library, need not be included in conveying the object code work.

A “User Product” is either (1) a “consumer product”, which means any tangible personal property which is normally used for personal, family, or household purposes, or (2) anything designed or sold for incorporation into a dwelling. In determining whether a product is a consumer product, doubtful cases shall be resolved in favor of coverage. For a particular product received by a particular user, “normally used” refers to a typical or common use of that class of product, regardless of the status of the particular user or of the way in which the particular user actually uses, or expects or is expected to use, the product. A product is a consumer product regardless of whether the product has substantial commercial, industrial or non-consumer uses, unless such uses represent the only significant mode of use of the product.

“Installation Information” for a User Product means any methods, procedures, authorization keys, or other information required to install and execute modified versions of a covered work in that User Product from a modified version of its Corresponding Source. The information must suffice to ensure that the continued functioning of the modified object code is in no case prevented or interfered with solely because modification has been made.

If you convey an object code work under this section in, or with, or specifically for use in, a User Product, and the conveying occurs as part of a transaction in which the right of possession and use of the User Product is transferred to the recipient in perpetuity or for a fixed term (regardless of how the transaction is characterized), the Corresponding Source conveyed under this section must be accompanied by the Installation Information. But this requirement does not apply if neither you nor any third party retains the ability to install modified object code on the User Product (for example, the work has been installed in ROM).

The requirement to provide Installation Information does not include a requirement to continue to provide support service, warranty, or updates for a work that has been modified or installed by the recipient, or for the User Product in which it has been modified or installed. Access to a network may be denied when the modification itself materially and adversely affects the operation of the network or violates the rules and protocols for communication across the network.

Corresponding Source conveyed, and Installation Information provided, in accord with this section must be in a format that is publicly documented (and with an implementation available to the public in source code form), and must require no special password or key for unpacking, reading or copying.

7. Additional Terms.

“Additional permissions” are terms that supplement the terms of this License by making exceptions from one or more of its conditions. Additional permis-

sions that are applicable to the entire Program shall be treated as though they were included in this License, to the extent that they are valid under applicable law. If additional permissions apply only to part of the Program, that part may be used separately under those permissions, but the entire Program remains governed by this License without regard to the additional permissions.

When you convey a copy of a covered work, you may at your option remove any additional permissions from that copy, or from any part of it. (Additional permissions may be written to require their own removal in certain cases when you modify the work.) You may place additional permissions on material, added by you to a covered work, for which you have or can give appropriate copyright permission.

Notwithstanding any other provision of this License, for material you add to a covered work, you may (if authorized by the copyright holders of that material) supplement the terms of this License with terms:

- a) Disclaiming warranty or limiting liability differently from the terms of sections 15 and 16 of this License; or
- b) Requiring preservation of specified reasonable legal notices or author attributions in that material or in the Appropriate Legal Notices displayed by works containing it; or
- c) Prohibiting misrepresentation of the origin of that material, or requiring that modified versions of such material be marked in reasonable ways as different from the original version; or

- d) Limiting the use for publicity purposes of names of licensors or authors of the material; or
- e) Declining to grant rights under trademark law for use of some trade names, trademarks, or service marks; or
- f) Requiring indemnification of licensors and authors of that material by anyone who conveys the material (or modified versions of it) with contractual assumptions of liability to the recipient, for any liability that these contractual assumptions directly impose on those licensors and authors.

All other non-permissive additional terms are considered “further restrictions” within the meaning of section 10. If the Program as you received it, or any part of it, contains a notice stating that it is governed by this License along with a term that is a further restriction, you may remove that term. If a license document contains a further restriction but permits relicensing or conveying under this License, you may add to a covered work material governed by the terms of that license document, provided that the further restriction does not survive such relicensing or conveying.

If you add terms to a covered work in accord with this section, you must place, in the relevant source files, a statement of the additional terms that apply to those files, or a notice indicating where to find the applicable terms.

Additional terms, permissive or non-permissive, may be stated in the form of a separately written license, or stated as exceptions; the above requirements apply either way.

8. Termination.

You may not propagate or modify a covered work except as expressly provided under this License. Any attempt otherwise to propagate or modify it is void, and will automatically terminate your rights under this License (including any patent licenses granted under the third paragraph of section 11).

However, if you cease all violation of this License, then your license from a particular copyright holder is reinstated (a) provisionally, unless and until the copyright holder explicitly and finally terminates your license, and (b) permanently, if the copyright holder fails to notify you of the violation by some reasonable means prior to 60 days after the cessation.

Moreover, your license from a particular copyright holder is reinstated permanently if the copyright holder notifies you of the violation by some reasonable means, this is the first time you have received notice of violation of this License (for any work) from that copyright holder, and you cure the violation prior to 30 days after your receipt of the notice.

Termination of your rights under this section does not terminate the licenses of parties who have received copies or rights from you under this License. If your rights have been terminated and not permanently reinstated, you do not qualify to receive new licenses for the same material under section 10.

9. Acceptance Not Required for Having Copies.

You are not required to accept this License in order to receive or run a copy of the Program. Ancillary propagation of a covered work occurring solely as

a consequence of using peer-to-peer transmission to receive a copy likewise does not require acceptance. However, nothing other than this License grants you permission to propagate or modify any covered work. These actions infringe copyright if you do not accept this License. Therefore, by modifying or propagating a covered work, you indicate your acceptance of this License to do so.

10. Automatic Licensing of Downstream Recipients.

Each time you convey a covered work, the recipient automatically receives a license from the original licensors, to run, modify and propagate that work, subject to this License. You are not responsible for enforcing compliance by third parties with this License.

An “entity transaction” is a transaction transferring control of an organization, or substantially all assets of one, or subdividing an organization, or merging organizations. If propagation of a covered work results from an entity transaction, each party to that transaction who receives a copy of the work also receives whatever licenses to the work the party’s predecessor in interest had or could give under the previous paragraph, plus a right to possession of the Corresponding Source of the work from the predecessor in interest, if the predecessor has it or can get it with reasonable efforts.

You may not impose any further restrictions on the exercise of the rights granted or affirmed under this License. For example, you may not impose a license fee, royalty, or other charge for exercise of rights granted under this License, and you may not initiate litigation

(including a cross-claim or counterclaim in a lawsuit) alleging that any patent claim is infringed by making, using, selling, offering for sale, or importing the Program or any portion of it.

11. Patents.

A “contributor” is a copyright holder who authorizes use under this License of the Program or a work on which the Program is based. The work thus licensed is called the contributor’s “contributor version”.

A contributor’s “essential patent claims” are all patent claims owned or controlled by the contributor, whether already acquired or hereafter acquired, that would be infringed by some manner, permitted by this License, of making, using, or selling its contributor version, but do not include claims that would be infringed only as a consequence of further modification of the contributor version. For purposes of this definition, “control” includes the right to grant patent sublicenses in a manner consistent with the requirements of this License.

Each contributor grants you a non-exclusive, worldwide, royalty-free patent license under the contributor’s essential patent claims, to make, use, sell, offer for sale, import and otherwise run, modify and propagate the contents of its contributor version.

In the following three paragraphs, a “patent license” is any express agreement or commitment, however denominated, not to enforce a patent (such as an express permission to practice a patent or covenant not to sue for patent infringement). To “grant” such a patent

license to a party means to make such an agreement or commitment not to enforce a patent against the party.

If you convey a covered work, knowingly relying on a patent license, and the Corresponding Source of the work is not available for anyone to copy, free of charge and under the terms of this License, through a publicly available network server or other readily accessible means, then you must either (1) cause the Corresponding Source to be so available, or (2) arrange to deprive yourself of the benefit of the patent license for this particular work, or (3) arrange, in a manner consistent with the requirements of this License, to extend the patent license to downstream recipients. “Knowingly relying” means you have actual knowledge that, but for the patent license, your conveying the covered work in a country, or your recipient’s use of the covered work in a country, would infringe one or more identifiable patents in that country that you have reason to believe are valid.

If, pursuant to or in connection with a single transaction or arrangement, you convey, or propagate by procuring conveyance of, a covered work, and grant a patent license to some of the parties receiving the covered work authorizing them to use, propagate, modify or convey a specific copy of the covered work, then the patent license you grant is automatically extended to all recipients of the covered work and works based on it.

A patent license is “discriminatory” if it does not include within the scope of its coverage, prohibits the exercise of, or is conditioned on the non-exercise of one or more of the rights that are specifically granted un-

der this License. You may not convey a covered work if you are a party to an arrangement with a third party that is in the business of distributing software, under which you make payment to the third party based on the extent of your activity of conveying the work, and under which the third party grants, to any of the parties who would receive the covered work from you, a discriminatory patent license (a) in connection with copies of the covered work conveyed by you (or copies made from those copies), or (b) primarily for and in connection with specific products or compilations that contain the covered work, unless you entered into that arrangement, or that patent license was granted, prior to 28 March 2007.

Nothing in this License shall be construed as excluding or limiting any implied license or other defenses to infringement that may otherwise be available to you under applicable patent law.

12. No Surrender of Others' Freedom.

If conditions are imposed on you (whether by court order, agreement or otherwise) that contradict the conditions of this License, they do not excuse you from the conditions of this License. If you cannot convey a covered work so as to satisfy simultaneously your obligations under this License and any other pertinent obligations, then as a consequence you may not convey it at all. For example, if you agree to terms that obligate you to collect a royalty for further conveying from those to whom you convey the Program, the only way you could satisfy both those terms and this License would be to refrain entirely from conveying the Program.

13. Use with the GNU Affero General Public License.

Notwithstanding any other provision of this License, you have permission to link or combine any covered work with a work licensed under version 3 of the GNU Affero General Public License into a single combined work, and to convey the resulting work. The terms of this License will continue to apply to the part which is the covered work, but the special requirements of the GNU Affero General Public License, section 13, concerning interaction through a network will apply to the combination as such.

14. Revised Versions of this License.

The Free Software Foundation may publish revised and/or new versions of the GNU General Public License from time to time. Such new versions will be similar in spirit to the present version, but may differ in detail to address new problems or concerns.

Each version is given a distinguishing version number. If the Program specifies that a certain numbered version of the GNU General Public License “or any later version” applies to it, you have the option of following the terms and conditions either of that numbered version or of any later version published by the Free Software Foundation. If the Program does not specify a version number of the GNU General Public License, you may choose any version ever published by the Free Software Foundation.

If the Program specifies that a proxy can decide which future versions of the GNU General Public License can be used, that proxy’s public statement of acceptance of

a version permanently authorizes you to choose that version for the Program.

Later license versions may give you additional or different permissions. However, no additional obligations are imposed on any author or copyright holder as a result of your choosing to follow a later version.

15. Disclaimer of Warranty.

THERE IS NO WARRANTY FOR THE PROGRAM, TO THE EXTENT PERMITTED BY APPLICABLE LAW. EXCEPT WHEN OTHERWISE STATED IN WRITING THE COPYRIGHT HOLDERS AND/OR OTHER PARTIES PROVIDE THE PROGRAM "AS IS" WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE. THE ENTIRE RISK AS TO THE QUALITY AND PERFORMANCE OF THE PROGRAM IS WITH YOU. SHOULD THE PROGRAM PROVE DEFECTIVE, YOU ASSUME THE COST OF ALL NECESSARY SERVICING, REPAIR OR CORRECTION.

16. Limitation of Liability.

IN NO EVENT UNLESS REQUIRED BY APPLICABLE LAW OR AGREED TO IN WRITING WILL ANY COPYRIGHT HOLDER, OR ANY OTHER PARTY WHO MODIFIES AND/OR CONVEYS THE PROGRAM AS PERMITTED ABOVE, BE LIABLE TO YOU FOR DAMAGES, INCLUDING ANY GENERAL, SPECIAL, INCIDENTAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF THE USE OR INABILITY TO USE THE PROGRAM (INCLUDING BUT NOT

LIMITED TO LOSS OF DATA OR DATA BEING RENDERED INACCURATE OR LOSSES SUSTAINED BY YOU OR THIRD PARTIES OR A FAILURE OF THE PROGRAM TO OPERATE WITH ANY OTHER PROGRAMS), EVEN IF SUCH HOLDER OR OTHER PARTY HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

17. Interpretation of Sections 15 and 16.

If the disclaimer of warranty and limitation of liability provided above cannot be given local legal effect according to their terms, reviewing courts shall apply local law that most closely approximates an absolute waiver of all civil liability in connection with the Program, unless a warranty or assumption of liability accompanies a copy of the Program in return for a fee.

END OF TERMS AND CONDITIONS

How to Apply These Terms to Your New Programs

If you develop a new program, and you want it to be of the greatest possible use to the public, the best way to achieve this is to make it free software which everyone can redistribute and change under these terms.

To do so, attach the following notices to the program. It is safest to attach them to the start of each source file to most effectively state the exclusion of warranty; and each file should have at least the “copyright” line and a pointer to where the full notice is found.

<one line to give the program's name and a brief idea of what it does>

Copyright (C) <textyear> <name of author>

B.2 TERMS AND CONDITIONS

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <http://www.gnu.org/licenses/>.

Also add information on how to contact you by electronic and paper mail.

If the program does terminal interaction, make it output a short notice like this when it starts in an interactive mode:

```
<program> Copyright (C) <year> <name of author>
```

```
This program comes with ABSOLUTELY NO WARRANTY; for details type 'show w'
This is free software, and you are welcome to redistribute it
under certain conditions; type 'show c' for details.
```

The hypothetical commands `show w` and `show c` should show the appropriate parts of the General Public License. Of course, your program's commands might be different; for a GUI interface, you would use an “about box”.

You should also get your employer (if you work as a programmer) or school, if any, to sign a “copyright disclaimer” for the program, if necessary. For more information on this, and how to apply and follow the GNU GPL, see <http://www.gnu.org/licenses/>.

The GNU General Public License does not permit incorporating your program into proprietary programs. If your program is a subroutine library, you may consider it more useful to permit linking proprietary applications with the library. If this is what you want to do, use the GNU Lesser General Public License instead of this License. But first, please read <http://www.gnu.org/philosophy/why-not-lgpl.html>.